

CONSUMPTION

JYOTIRMOY BHATTACHARYA

1. TWO-PERIOD CASE

1.1. Budget constraint. Consider a consumer who lives for two periods, has an endowment of y_1 and y_2 units of goods in the two periods respectively and can borrow and lend any amount that they like at the real rate of interest r .

Suppose the consumer consumes c_1 in the first period. Then she will have to take a loan of $c_1 - y_1$ to finance her consumption. (This number can be negative, in which case the consumer is lending rather than borrowing.) In the next period the consumer will therefore have to make loan repayments of $(1 + r)(c_1 - y_1)$. Assume that the consumer does not want to make any bequests and cannot die with any outstanding loans, consumption in the second period must be,

$$c_2 = y_2 - (1 + r)(c_1 - y_1)$$

Simplifying and rearranging we have

$$c_1 + \frac{c_2}{1 + r} = y_1 + \frac{y_2}{1 + r} \quad (1)$$

This is the budget constraint faced by the consumer. We can interpret this to mean that the present value of the consumer's consumption stream must equal the present value of their incomes.

1.2. Utility maximization. Suppose the consumer maximises a quasiconcave utility function $U(c_1, c_2)$ subject to this budget constraint. Then the consumer's first-order conditions are

$$U_1(c_1, c_2) = \lambda \quad (2)$$

$$U_2(c_1, c_2) = \lambda / (1 + r) \quad (3)$$

where λ is the Lagrange multiplier corresponding to the budget constraint and $U_i(c_1, c_2)$ denotes the partial derivative $\partial U / \partial c_i$. We have explicitly shown the dependence of the partial derivatives on the value of consumption in both periods. These first-order conditions

along with the budget constraint (1) together determines the value of c_1 , c_2 and λ .

1.3. Comparative statics. Assuming that consumption in both periods is a normal good, an increase in either y_1 or y_2 increases both c_1 and c_2 .

The effects of a change in r are ambiguous. An increase in r makes consumption in period 2 relatively cheap compared to consumption in period 1. Therefore the substitution effect causes c_1 to decrease and c_2 to increase. It is traditional to decompose the income effect into two parts. First, an increase in r reduces the present value of the consumer's endowments and hence decreases his real income. Second, an increase in r , by making the consumption in period 2 cheaper increases his real income.¹ The sign of the resultant of these two effects on consumption depends on whether the consumer is a net lender in period 1 and a net borrower in period 2 or vice-versa. In case the consumer is a net lender in period 1 and a net borrower in period 2 the net income effect is positive. Assuming the consumption in both periods in a normal good, this means that the substitution effect and the income effect act in opposite directions on c_1 in this case leading to an ambiguous effect.

2. MANY PERIODS

Assume that rather than just living for two periods the consumer lives for $T + 1$ periods. Further assume that the real rate of interest takes a constant value r over the consumer's lifetime. For convenience we define $\delta = 1/(1 + r)$. It is also convenient to start time from period 0 rather than period 1.

2.1. Budget constraint. Arguing as before, the consumer's budget constraint is

$$\sum_{i=0}^T \delta^i c_i = \sum_{i=0}^T \delta^i y_i \quad (4)$$

2.2. Utility function. We could proceed as before by assuming a utility function $U(c_0, \dots, c_T)$ and deriving the first order conditions. However, because the marginal utility in each period depends on consumption in all periods it is hard to draw any sharp conclusions at this level of generality. Therefore we need to impose some restrictions on the form of the utility functions.

¹For more about the Slutsky equation in the case of a consumer with fixed endowments of goods see section 9.1 in Varian's *Microeconomic Analysis*, 3rd ed.

Suppose, for example we assume that the utility function is additively separable, i.e.

$$U(c_0, \dots, c_T) = v_0(c_0) + v_1(c_1) + \dots + v_T(c_T) \quad (5)$$

Then the first-order conditions take the form

$$v'_i(c_i) = \delta^i \lambda \quad i = 0, \dots, T \quad (6)$$

where, as before, λ is the Lagrange multiplier corresponding to the budget constraint.

Sometimes we want to restrict the consumers preferences even further, by assuming that the different v_i differ from each other by only a geometric discounting factor.

$$U(c_0, \dots, c_T) = \sum_{i=0}^T \beta^i u(c_i) \quad (7)$$

where β is a constant, referred to as the subjective rate of discount, such that $0 < \beta < 1$.

In this case the first-order conditions take the particularly simple form

$$u'(c_i) = \left(\frac{\delta}{\beta}\right)^i \lambda \quad i = 0, \dots, T \quad (8)$$

In case $\delta = \beta$, this implies that $u'(c_i)$ is the same for all i , which, assuming that $u'(\cdot)$ is a strictly decreasing function, means that c_i is constant for all i . The present period's income does not influence the present period's consumption at all. Consumption is determined solely by lifetime resources as given by (4).

The case $\delta \neq \beta$ is also instructive. Suppose $\delta > \beta$. In this case it follows from (8) that consumption decreases over time. Formally, this is because if $\delta > \beta$ then by (8) $u'(c_i)$ increases over time, and since $u'(c)$ is a decreasing function of consumption, this implies that c decreases over time.

The economic logic behind this result is that δ is the number of units of consumption we have to give up at present in order to purchase one more unit of consumption next period, whereas β is the number of units of marginal utility we are willing to give up at present in order to have one more unit of marginal utility in the next period. Suppose we start with the same consumption c in this period and the next. If we reduce consumption in the next period by a small amount Δc then at the prevailing market prices we can increase present consumption by $\delta \Delta c$. The increase in utility from

the increase in present consumption is approximately $u'(c)(\delta\Delta c)$.² The decrease in utility from the reduction in next period's consumption is approximately $\beta u'(c)(\Delta c)$. The net change in utility would be $(\delta - \beta)u'(c)(\Delta c)$ which is positive when $\delta > \beta$. Thus it is beneficial to increase present consumption and reduce future consumption if we are starting from a position of equality. Indeed, it will be optimal to increase consumption in the present period (say period i) and decrease consumption in the next period (period $i + 1$) till the following equality between the MRS and the price ratio is satisfied,

$$\frac{u'(c_{i+1})}{u'(c_i)} = \frac{\delta}{\beta}$$

If δ/β is close to 1 then c_{i+1} is close to c_i and we can use Taylor's Theorem from calculus to the above equation to the above equation to get a useful approximation.

$$\begin{aligned}\frac{u'(c_{i+1})}{u'(c_i)} &= \frac{\delta}{\beta} \\ \frac{u'(c_{i+1}) - u'(c_i)}{u'(c_i)} &= \frac{\delta}{\beta} - 1\end{aligned}$$

Applying Taylor's Theorem

$$\frac{u''(c_i)(c_{i+1} - c_i)}{u'(c_i)} \approx \frac{\delta}{\beta} - 1$$

Defining $\Delta c = c_{i+1} - c_i$, and dropping the subscript i ,

$$\left(\frac{u''(c)c}{u'(c)}\right) \left(\frac{\Delta c}{c}\right) \approx \frac{\delta}{\beta} - 1$$

The quantity $\sigma = -u''(c)c/u'(c)$ is known as the *intertemporal elasticity of substitution* and captures the sensitivity of marginal utility of changes in consumption. It is positive since marginal utility decreases with consumption.

$$\left(\frac{\Delta c}{c}\right) \approx \frac{1}{\sigma} \left(1 - \frac{\delta}{\beta}\right)$$

The formula confirms our earlier reasoning that consumption decreases over time if $\delta > \beta$. Moreover, it shows that the sensitivity of the growth of consumption on the rate of return depends inversely on the intertemporal elasticity of substitution. The reason for this is

²We are using Taylor's theorem: $u(c + \delta\Delta c) - u(c) \approx u'(c)(\delta\Delta c)$

that the more elastic is marginal utility to consumption, the smaller is the deviation in consumption from a constant path that is required to equate the ratio of marginal utilities in consecutive time periods to δ/β .

2.3. Exogenous variables. It is possible to unify (5) and (7) by writing

$$v_i(c_i) = \beta^i u(c_i, \xi_i)$$

where ξ_i is an exogenous variable such as the consumer's age or the number of members in the household. In this case the first-order conditions become

$$u'(c_i, \xi_i) = \left(\frac{\delta}{\beta}\right)^i \lambda \quad i = 0, \dots, T$$

Knowing how ξ affects the marginal utility would now let us make some predictions regarding the path of consumption.

2.4. Comparative statics. Assuming that consumption in every period is a normal good, an increase in y_i increases every c_i .

The effect of an increase in r , or equivalently, a decrease in δ remains ambiguous because of the same income and substitution effects as discussed earlier. But for the utility function given by (7), we can say a little more. From (8) we can see that a decrease in δ means that the *growth rate* of consumption speeds up. Remember that even in this case we do not have any information regarding the *level* of consumption in any period since the level would depend on λ which in turn depends on δ .

3. DYNAMIC PROGRAMMING

3.1. The setup. In a dynamic optimisation problem, our goal is to find a *path* of the choice variable which maximises the value of an objective function defined over the entire path of the choice variable. Often, there are constraints on what paths can be chosen. For example, in the consumption-saving problem we choose a path of consumption which maximises the lifetime utility function subject to a budget constraint.

The dynamic programming approach to solving dynamic optimisation problems turns this single large optimization problem into a sequence of simple optimization problems. At each point of time we try to find the best action at that particular point of time. But since this is a dynamic problem after all, this search for the best actions at a particular point of time has to be done with an eye on both the past

and the future. Past events and actions³ determine what choices can be made now. By the same token, the action that we take now will change the options available to us in the future. The value of the objective function that will be achieved will in general depend on the entire path of past, present and future actions and not just the action in any period in isolation.

In the dynamic programming framework this linkage between the past and the future is captured by the notion of the *state*. Intuitively, we can think of the state at any given point of time as a description of all the *relevant* information about the actions and events that have happened until that point. The state should contain all the information that is required from the decision-maker's history to determine the set of available actions at future points of time and to evaluate the contribution made by future actions to the objective function.

The notion that knowing the state at a point of time is enough to know what actions are possible in the future is captured by the following definitions:

Set of states (S_t): this is the set of possible states the decision-maker can be in time t . There is one such set for each time period t . The elements of these sets, i.e. the possible states, are assumed to be vectors with real-number elements. Elements of this set are denoted by s_t .

Set of actions (A_t): this is the set of possible actions that can be taken at time t . Elements of this set are denoted by a_t . As we shall see next, all possible actions cannot be taken at all possible states.

Constraint correspondence $f_t(s_t) \subset A_t$: this tells us the subset of actions that are available in a particular state. This is not a function but a correspondence (i.e. a set-valued function) since for each element of S_t it gives us a subset and not just a single element of A_t .

Transition function $\Gamma_t(s_t, a_t) \in S_{t+1}$: This tells us our state in period $t + 1$ if we take the action a_t in state s_t in period t .

With these definitions in hand we can define the set of *feasible plans* when starting with s_t at time t , denoted by $\Phi_t(s_t)$, as the set of sequences of actions (a_t, \dots, a_T) such that

$$a_i \in f_i(s_i) \quad \text{for } i = t, \dots, T$$

³We want to make a distinction between *events* which are outside of our control and *actions* which are things we choose. This distinction becomes important when we are dealing with uncertainty.

and

$$s_{i+1} = \Gamma_i(s_i, a_i) \quad \text{for } i = t, \dots, T-1$$

The first condition says that the action taken on each date is a feasible action given the state. The second condition says that the state at each date is derived from the state and action taken in the previous date, with the state at time t as given.

The set of feasible plans tells us about the constraints faced in our optimisation problem. What about the objective function? We assume that the objective function can be written in an additively separable form

$$U_t(a_t, s_t, \dots, a_T, s_T) = \sum_{i=t}^T v_i(a_i, s_i)$$

where $v_t(a_t, s_t)$ is the *per-period payoff function* that gives the contribution of action a_t in state s_t at time t to the overall objective. Being able to write the objective function in an additively separable form is essential for us to be able to use dynamic programming.⁴

In writing the above objective function we have also assumed that there is a finite time period T at which our optimisation problem comes to an end. This assumption of what is known as a finite horizon is made just to simplify the mathematics. Dynamic programming problems with an infinite horizon are routinely used in economic modelling.

The solution to the dynamic programming problem is expressed in terms of two functions:

Policy function $g_t(s_t) \in f_t(s_t)$: The policy function tells us the best action to take in each possible state at time t among all the available actions. In general it is possible that there be two equally good actions at a particular state, in which case the policy function would have to be replaced by the policy correspondence.

Value function $V_t(s_t) \in \mathbb{R}$: The value function denotes the maximum attainable value of the objective function when starting at time t from state s_t . That is,

$$\begin{aligned} V_t(s_t) &= \max_{(a_t, \dots, a_T) \in \Phi_t(s_t)} U_t(a_t, s_t, \dots, a_T, s_T) \\ &= \max_{(a_t, \dots, a_T) \in \Phi_t(s_t)} [v_t(a_t, s_t) + \dots + v_T(a_T, s_T)] \end{aligned}$$

⁴We are cheating a bit here. The assumption of additive separability can be relaxed to what is called ‘recursiveness’ while still allowing the use of dynamic programming.

In applications of dynamic programming we generally want to know the optimal path starting at a specific point of time (taken to be $t = 0$ here) and from a particular state at that point of time (say \bar{s}_0). But we have defined the policy and value functions for all points of time and for each possible state at each of the time periods. Thus it would seem that we have multiplied our work manifold beyond what is necessary for our original problem. But as we shall see below, being willing to contemplate the policy and value functions for all possible time periods and states often actually simplifies the task of solving the original problem.

3.2. Bellman's Principle of Optimality. Suppose I am starting at some time $t < T$ from some particular state s_t and trying to find the best actions from time t to T , where 'best' means the choices of actions and consequent states which maximise

$$U_t(a_t, s_t, \dots, a_T, s_T) = \sum_{i=t}^T v_i(a_i, s_i).$$

Because of the additive nature of the lifetime utility function we can rewrite the above equation as

$$U_t(a_t, s_t, \dots, a_T, s_T) = v_t(a_t, s_t) + U_{t+1}(a_{t+1}, s_{t+1}, \dots, a_T, s_T)$$

If we divide the plan (path of actions and corresponding states) from time t to time T into a "head" consisting of the action in period t and a "tail" consisting of actions in period $t + 1$ to T then the above equation says that the lifetime utility of the plan starting at period t is the sum of the per-period payoff at time t (the value of the "head") and the lifetime utility of the remaining part of the plan from period $t + 1$ onwards (the value of the "tail").

How do we find the plan which maximises U_t ? Suppose we choose the action \hat{a}_t in period t . This will lead us to the state $\hat{s}_{t+1} = \Gamma_t(s_t, \hat{a}_t)$ in the next period. Now we have to pick a plan from period $t + 1$ onward. Now $U_t = v_t(\hat{a}_t, s_t) + U_{t+1}$ and $v_t(\hat{a}_t, s_t)$ is already fixed by our choice of action \hat{a}_t in period t . Therefore in choosing our plan from period $t + 1$ onward the best we can do is to pick a plan that maximises U_{t+1} . This optimal plan for the "tail" yields the value of U_{t+1} equal to $V_{t+1}(\hat{s}_{t+1})$. Thus we can evaluate each choice of action \hat{a}_t in the "head" by looking at

$$\tilde{U}_t = v_t(\hat{a}_t, s_t) + V_{t+1}(\hat{s}_{t+1}), \quad \text{where } \hat{s}_{t+1} = \Gamma(s_t, \hat{a}_t)$$

We have put a tilde over U_t to remind ourselves that now we are not considering arbitrary plans starting at t but only plans where the

“tail” component is optimal given the state \hat{s}_{t+1} at which we find ourselves in the beginning of period $t + 1$.

The optimal plan from period t involves choosing \hat{a}_t which maximises the expression above. Since the value function gives the value of the objective function U_t for the optimal plan, it is therefore the case that,

$$V_t(s_t) = \max_{\hat{a}_t \in f_t(s_t)} [v_t(\hat{a}_t, s_t) + V_{t+1}(\Gamma(s_t, \hat{a}_t))], \quad \text{for } t < T \quad (9)$$

Equation (9) above which relates the value function at consecutive time periods is known as *Bellman's Equation*. The argument above, which shows that the value function must satisfy Bellman's equation is known as *Bellman's Principle of Optimality*.⁵

Intuitively Bellman's equation tells us that we can evaluate each present action by adding its contribution $v_t(\hat{a}_t, s_t)$ to the objective in the present period and the value $V_{t+1}(\hat{s}_{t+1})$ of the state in which it leaves us in the next period. Provided we know the function V_{t+1} for all possible states in the next period we can choose the best action in the current period by choosing \hat{a}_t to maximise this sum. Thus we have turned the big optimisation problem of choosing an entire sequence of actions from time 0 to time T into a sequence of simple optimisation problems, one for each time period t , in each of which we choose a single action \hat{a}_t .

But there seems to be a chicken-and-egg problem: we cannot use Bellman's equation without knowing V_{t+1} for each t and how do we know V_{t+1} if we have not solved the optimisation problem already? Here our finite horizon assumption makes life particularly simple for us.

Since period T is the last period, our objective function in that period is

$$U_T(a_T, s_T) = v_T(a_T, s_T)$$

and the value function is simply given by

$$V_T(s_T) = \max_{a_T \in f_T(s_T)} v_T(a_T, s_T)$$

We can solve this maximisation problem and calculate V_T since $v_T(\cdot, \cdot)$ is a known function.

⁵To be complete, Bellman's principle of optimality also deals with the converse: that a function which satisfies Bellman's equation plus some other technical conditions must be the value function. This converse is not important in our current finite horizon setting.

Now consider Bellman's equation for period $T - 1$:

$$V_{T-1}(s_{T-1}) = \max_{\hat{a}_{T-1} \in f_t(s_{T-1})} [v_{T-1}(\hat{a}_{T-1}, s_{T-1}) + V_T(\Gamma(s_{T-1}, \hat{a}_{T-1}))]$$

As we have calculated $V_T(\cdot)$ in the previous step, all the functions in the maximisation problem are known and we can solve the problem to calculate V_{T-1} . With this in hand we can solve Bellman's equation for period $T - 2$. We keep going backward one period at a time until we have calculated the value function for all periods until period 0. At each step the value of \hat{a}_t , as a function of s_t , which solves the maximisation problem gives us the policy function. So by the end of our process we also have the policy function for each time period.

Now if we are given a starting state \bar{s}_0 in period 0 we can use the calculated policy function for period 0 to find the best action a_0 in period 0. We know from the transition function that we will end up in state $s_1 = \Gamma(\bar{s}_0, a_0)$ in the next period. The policy function for period 1 tells us the best action a_1 to take in that period. We again use the transition function to tell us the next state $s_2 = \Gamma(s_1, a_1)$. And so on until we have traced out the optimal plan to time T . Our optimisation problem is solved!

The way we have calculated the value function backwards from a known final time period is sometimes called "backward induction".

3.3. Example: consumption-savings with log utility. Suppose that the consumer maximises

$$\sum_{i=0}^T \log(c_i)$$

subject to

$$\sum_{i=0}^T c_i / R^i = w_0$$

Can we solve this maximisation problem using dynamic programming? The action variable in this case must be c_i since it is the variable being chosen by the decision maker. The objective function is already in an additively separable form with a per-period payoff $\log(c_i)$. But what is the state?

Since the per-period payoff depends only on the action variable c_i we do not need any notion of state to evaluate payoffs. But the choice of a consumption in each period does affect future periods through the budget. The more we consume today, the less purchasing power

we have to consume tomorrow. We can capture this by rearranging the budget slightly to read,

$$\sum_{i=1}^T c_i / R^i = w_0 - c_0$$

Multiplying throughout by R we have

$$\sum_{i=1}^T c_i / R^{(i-1)} = R(w_0 - c_0)$$

Which shows us that the path of consumption from time 1 onwards follows a budget constraint of the same form as the period 0 budget constraint provided we take

$$w_1 = R(w_0 - c_0)$$

This suggests to us that we can take the wealth at the beginning of period t as our state with the transition function,

$$w_{t+1} = R(w_t - c_t), \quad \text{for } t = 0, \dots, T$$

and the constraint function

$$c_T = w_T$$

The constraint captures the fact that there can be no outstanding debt in the last period and a consumer who has monotonic preferences would not leave any wealth unused in the last period. There is no constraint on consumption in periods other than T .⁶

You can check that we have formulated the problem right by eliminating w_1, \dots, w_T in the transition functions and constraints above to recover our original budget constraint.

Now we can use backward induction to calculate the value function and policy function for all time periods. Denoting the policy function by $g(\cdot)$ and the value function in period t by $V_t(\cdot)$, we have,

$$g_T(w) = w, \quad V_T(w_T) = \log(g_T(w)) = \log(w) \quad (10)$$

Now consider period $T - 1$. Bellman's principle of optimality tells us,

$$\begin{aligned} V_{T-1}(w) &= \max_c [\log(c) + V_T(R(w - c))] \\ &= \max_c [\log(c) + \log(R(w - c))] \quad [\text{using (10)}] \end{aligned} \quad (11)$$

⁶We could have imposed a non-negativity constraint but we leave it out for simplicity

The first-order condition for this maximisation problem is:

$$\begin{aligned}\frac{1}{c} + \frac{-R}{R(w-c)} &= 0 \\ w - c &= c \\ c &= w/2\end{aligned}$$

Since the objective function in (11) is concave in c (check this!), the first-order condition is sufficient and gives us our policy function:

$$g_{T-1}(w) = w/2$$

Substituting this into (11) we get the value function,

$$\begin{aligned}V_{T-1}(w) &= \log(g_{T-1}(w)) + V_T(R(w - g_{T-1}(w))) \\ &= \log(w/2) + \log(Rw/2) \\ &= \log(R) + 2\log(w/2)\end{aligned}\tag{12}$$

Now that we know V_{T-1} we could use the Bellman equation relating V_{t-2} to V_{t-1} to derive g_{T-2} and V_{T-2} . If you do this you will find,

$$g_{T-2}(w) = w/3, \quad V_{T-2}(w) = (1+2)\log(R) + 3\log(w/3) \tag{13}$$

We could continue like this to find V_{T-3}, \dots, V_0 . In general this is precisely what we do. In fact, in most applications of dynamic programming it is not possible to express the value function by a formula in the state variables and the best that we can do is to use a computer to calculate the value function at a number of possible values of the state variable using Bellman's equation.

But our present problem is a particularly simple one. Looking at (12) and (13) suggests to us the guess,

$$V_{T-n}(w) = \frac{n(n+1)}{2} \log(R) + (n+1) \log\left(\frac{w}{n+1}\right) \tag{14}$$

[Remember $1 + 2 + \dots + n = n(n+1)/2$]

How do we check that our guess is right? We will use the principle of mathematical induction. By comparing to (12) we see that (14) is correct for $n = 1$. Suppose that the equation is true for $n = k$. What then would be $V_{T-(k+1)}$? We once again write down the Bellman equation

$$\begin{aligned}
V_{T-(k+1)} &= \max_c [\log(c) + V_{T-k}(R(w-c))] \\
&= \max_c \left[\log(c) + \frac{k(k+1)}{2} \log(R) + (k+1) \log \left(\frac{R(w-c)}{k+1} \right) \right] \\
&\quad [\text{assuming (14)}]
\end{aligned} \tag{15}$$

The first-order condition is:

$$\begin{aligned}
\frac{1}{c} + (k+1) \left(\frac{k+1}{R(w-c)} \right) \left(\frac{-R}{k+1} \right) &= 0 \\
(k+1) \frac{1}{(w-c)} &= \frac{1}{c} \\
c &= w/(k+2)
\end{aligned} \tag{16}$$

Substituting this into (15) we have

$$\begin{aligned}
V_{T-(k+1)} &= \log(c) + \frac{k(k+1)}{2} \log(R) + (k+1) \log \left(\frac{R(w-c)}{k+1} \right) \\
&\quad \text{substituting (16),} \\
&= \log \left(\frac{w}{k+2} \right) + \frac{k(k+1)}{2} \log(R) + (k+1) \log \left(\frac{Rw}{k+2} \right) \\
&= \log \left(\frac{w}{k+2} \right) + \frac{k(k+1)}{2} \log(R) + (k+1) \log(R) + (k+1) \log \left(\frac{w}{k+2} \right) \\
&= \frac{(k+2)(k+1)}{2} \log(R) + (k+2) \log \left(\frac{w}{k+2} \right)
\end{aligned} \tag{17}$$

But this is the same as (14) for $n = k+1$. We therefore conclude that if (14) is true for $n = k$ it is also true for $n = k+1$. We have already checked that (14) is true for $n = 1$. Hence we conclude by the principle of mathematical induction that the value function for our dynamic programming problem is given by (14) for $n = 1, \dots, T$.

Also, now that we have verified that (14) is indeed the value function of the problem, (16) gives the policy function, i.e.

$$g_{T-n}(w) = w/(n+1) \tag{18}$$

4. THE EULER EQUATION

As we discussed in the last section, for most dynamic programming problems it is not possible to compute the value and policy

functions in terms of simple formulae. The best we can do is to calculate numerical values. But even if we cannot find an exact formula for the solution to our optimisation problem, it may still be possible to get some qualitative information about the problem by studying the consequences of the Bellman equation. That is the subject of this section.

Let's recall the Bellman equation,

$$V_t(w_t) = \max_{c_t} [u(c_t) + V_{t+1}(R(w_t - c_t))]$$

The first order condition for this maximisation problem is:

$$u'(c_t) = RV'_{t+1}(R(w_t - c_t)) \quad (19)$$

By itself (19) does not seem very useful unless we know $V_{t+1}(\cdot)$ and can calculate its derivative. But there is a trick that we can use to eliminate this unknown derivative from (19).⁷

Let $c_t^*(w_t)$ be the optimal consumption in period t when period t wealth is w_t . From (19) we already know that,

$$u'[c_t^*(w_t)] = RV'_{t+1}[R(w_t - c_t^*(w_t))] = RV'_{t+1}(w_{t+1}) \quad (20)$$

⁷The 'trick' is a particular case of a general result known as the envelope theorem. See section M.L of Mas-Colell, Whinston and Green or some mathematical methods book for more detail.

But from the definition of the value function

$$V_t(w_t) = u[c_t^*(w_t)] + V_{t+1}[R(w_t - c_t^*(w_t))]$$

Differentiating with respect to w_t we have,

$$\begin{aligned} V_t'(w_t) &= u'[c_t^*(w)]c_t^{*'}(w_t) + V_{t+1}'[R(w_t - c_t^*(w_t))][R(1 - c_t^{*'}(w_t))] \\ &= c_t^{*'}(w_t)[u'(\cdot)] - RV_{t+1}'(\cdot) + RV_{t+1}'[R(w_t - c_t^*(w_t))] \end{aligned}$$

From (20) the first terms equals 0, so,

$$V_t'(w_t) = RV_{t+1}'(w_{t+1})$$

Using (20)

$$V_t'(w_t) = u'(c_t)$$

The equation above was derived for arbitrary t . So it is equally good for $t + 1$, i.e.

$$V_{t+1}'(w_{t+1}) = u'(c_{t+1})$$

Substituting this in (20) we have,

$$\frac{u'(c_{t+1})}{u'(c_t)} = \frac{1}{R} = \delta \quad (21)$$

This condition is known as the Euler⁸ equation for our dynamic programming problem. We can alternatively derive it by starting out with an optimal consumption plan, increasing consumption in period t by a small amount Δc and reducing consumption in period $t + 1$ by $R\Delta c$ so that wealth at the end of the period $t + 1$ is once again the same as what it would have been under the optimal plan. The first-order change in utility from this deviation is

$$\Delta u = u'(c_t)[\Delta c] - u'(c_{t+1})[R\Delta c]$$

Now for the original plan to have been optimal Δu must be 0 since if $\Delta u > 0$ the deviation considered above increases total utility whereas if $\Delta u < 0$ then the opposite of the deviation considered above increases total utility. But $\Delta u = 0$ implies

$$u'(c_t) - Ru'(c_{t+1}) = 0$$

which is again our Euler equation (21).

The Euler equation also follows from the first-order conditions (6) of the Lagrange-multiplier approach, showing that we have come full circle.

⁸Pronounced "oiler". Named after a eighteenth-century mathematician who was among the earliest to study dynamic optimisation problems.

The Euler equation tells us how consumption should grow or decline. It does not tell us the level of the consumption. But we can characterise the entire consumption path if we keep track of the path of wealth implied by the path of consumption and impose, in addition to the Euler equation, the conditions

$$w_0 = \overline{w_0}$$

which comes to us as a given data and

$$w_T = 0$$

which comes to us from our no bequest, no terminal borrowing, monotonic utility assumptions about the terminal period.

5. UNCERTAINTY

5.1. Euler equation. In the case of uncertainty in labour incomes, but with certain interest rates, the Euler equation becomes

$$v'_t(c_t) = RE_t[v'_{t+1}(c_{t+1})] \quad (22)$$

where E_t denotes the mathematical expectation conditional on information at time t .

5.2. Quadratic felicity.

5.2.1. Martingale property. Suppose the felicity (i.e. per-period utility) function is

$$v_t(c_t) = \beta^t(ac_t - 0.5c_t^2)$$

where a is some constant.

In this case (22) specialises to

$$a - c_t = R\beta(a - E_t c_{t+1})$$

If we further assume that $R = 1/\beta$ then

$$E_t c_{t+1} = c_t \quad (23)$$

that is, consumption is a martingale process.

Since c_t is part of the information set at time t , $E_t c_t = c_t$. Therefore, another way to write (23) is

$$E_t(c_{t+1} - c_t) = 0$$

which says that the change in consumption between time t and $t + 1$ has no predictable direction based on information at time t .

This result is a consequence of the very special assumptions that we have made. Assuming the same felicity function for each period (apart from the discount factor β) and then assuming that the

market rate of discount ($1/R$) equals this subjective discount factor creates a situation where the consumer has no desire to have a higher consumption in any particular period of her life either to meet greater consumption needs or to take advantage of the difference between market and subjective discount rates. Unconstrained lending and borrowing mean that the consumer can actually move around her income across periods so as to achieve this perfect symmetry in her consumption in the sense of equating expected marginal utility across periods. But with quadratic felicity expected marginal utility is the same thing as expected consumption and we have our martingale result.

The long list of assumptions leading up to the martingale result means that this precise result is not very robust or realistic. Therefore rather than taking it as a property that is likely to be literally true, we should understand it as a demonstration of the tendency of the lending and borrowing behaviour of consumers to delink current consumption from current income. This tendency will be there as long as consumers have access to asset markets, though in more realistic settings it will be overlaid with factors which impart a systematic pattern to the trajectory of consumption such as a changing pattern of lifetime consumption needs or differences between the subjective and market rate of discount.

5.2.2. A martingale lemma. The definition of a martingale tells us about the expectation of a process at a period conditional on information available in the immediately preceding period. However, sometimes we need to find the expectation of a process at a period conditional on information available much further back in the past. The following lemma helps us.

Lemma 1. *If X_t is a martingale then for any m and any $n > 0$,*

$$E[X_{m+n}|m] = X_m.$$

Proof. The proof is by mathematical induction on n .

For $n = 1$ the proof follows directly from the definition of a martingale.

Suppose the lemma is true for $n = k$. Consider the case $n = k + 1$. Noting that the information available at time $m + k$ is a superset of the information available at time m , we have from the law of iterated expectations

$$E[X_{m+k+1}|m] = E[E[X_{m+k+1}|m+k]|m]$$

The martingale property, applied at time $m + k$ tells us that

$$E[X_{m+k+1}|m+k] = X_{m+k}$$

. The assumption that the lemma is true for $n = k$ gives us,

$$E[X_{m+k}|m] = X_m$$

Putting everything together, we have

$$E[X_{m+k+1}|m] = E[E[X_{m+k+1}|m+k]|m] = E[X_{m+k}|m] = X_m$$

thus establishing the result for $n = k + 1$.

Since we have shown that the result is true for $n = 1$ and it is true for $n = k + 1$ whenever it is true for $n = k$, it follows from the principle of mathematical induction that it is true for all $n > 0$. \square

5.2.3. The level of consumption. The martingale property of consumption only tells us how consumption evolves from one point to the next, not the *level* of consumption. The level of consumption would depend on the consumer's resources in terms of her initial wealth and expected labour income. We now show that this is so mathematically by deriving an explicit formula for the level of consumption in the case where consumption is a martingale.

Consider a consumer who stands at period t with wealth w_t and is planning her future consumption for the periods $t, t + 1, \dots, T$. Since she cannot leave any bequests or outstanding debt in period T , it must be the case that her *realized* stream of consumption (c_t) and labour income (y_t) must satisfy,

$$\sum_{i=t}^T \delta^{i-t} c_i = \sum_{i=t}^T \delta^{i-t} y_i + w_t \quad (24)$$

Taking expectations as of time t ,

$$\sum_{i=t}^T \delta^{i-t} E_t c_i = \sum_{i=t}^T \delta^{i-t} E_t y_i + w_t \quad (25)$$

From Lemma 1 on page 17, $E_t c_i = c_t$ for all $i > t$. And for $i = t$, $E_t c_t = c_t$ because c_t is known at time t . Hence,

$$c_t \sum_{i=t}^T \delta^{i-t} = \sum_{i=t}^T \delta^{i-t} E_t y_i + w_t \quad (26)$$

$$c_t = \frac{1}{\sum_{i=t}^T \delta^{i-t}} \left[\sum_{i=t}^T \delta^{i-t} E_t y_i + w_t \right] \quad (27)$$

Thus the level of consumption in a given period depends on the expected discounted value of the future stream of labour income over the entire remaining lifetime as well as initial wealth. This once again reiterates the idea of the permanent income hypothesis that the consumption in each period depends not just on income in that period but on the entire expected path of future income.

5.2.4. *Increments in consumption.* With an explicit formula for the level of consumption in hand, we can now try to understand the martingale result better by seeing what it is exactly that drives changes in consumption.

Rewriting (26) for period $t + 1$ we have,

$$c_{t+1} \sum_{i=t+1}^T \delta^{i-t-1} = \sum_{i=t+1}^T \delta^{i-t-1} E_{t+1} y_i + w_{t+1}$$

Substituting $w_{t+1} = (w_t + y_t - c_t) / \delta$,

$$c_{t+1} \sum_{i=t+1}^T \delta^{i-t-1} = \sum_{i=t+1}^T \delta^{i-t-1} E_{t+1} y_i + (w_t + y_t - c_t) / \delta$$

Multiplying throughout by δ ,

$$c_{t+1} \sum_{i=t+1}^T \delta^{i-t} = \sum_{i=t+1}^T \delta^{i-t} E_{t+1} y_i + w_t + y_t - c_t$$

Subtracting (26) from this equation

$$(c_{t+1} - c_t) \sum_{i=t+1}^T \delta^{i-t} - c_t = \sum_{i=t+1}^T \delta^{i-t} (E_{t+1} y_i - E_t y_i) + y_t - E_t y_t - c_t$$

But $E_t y_t = y_t$ since income at time t is known at that time,

$$\begin{aligned} (c_{t+1} - c_t) \sum_{i=t+1}^T \delta^{i-t} &= \sum_{i=t+1}^T \delta^{i-t} (E_{t+1} y_i - E_t y_i) \\ c_{t+1} - c_t &= \frac{1}{\sum_{i=t+1}^T \delta^{i-t}} \left[\sum_{i=t+1}^T \delta^{i-t} (E_{t+1} y_i - E_t y_i) \right] \end{aligned}$$

We can divide both the numerator and denominator by δ to have the factor multiplying the first term in the sums equal to one.⁹ This gives

⁹This is a purely aesthetic change and does not change any results.

us our final formula,

$$c_{t+1} - c_t = \frac{1}{\sum_{i=t+1}^T \delta^{i-t-1}} \left[\sum_{i=t+1}^T \delta^{i-t-1} (E_{t+1}y_i - E_t y_i) \right] \quad (28)$$

What the formula above says is that changes in consumption are a result of *revisions* in expectations of future income based on the difference in information at time t and $t + 1$. Therefore changes in income that were predictable at time t do not contribute to the change in consumption between time t and $t + 1$. This is consistent with our assumption that consumption is a martingale but goes further by predicting the actual size of the consumption change rather than just asserting that the expected value of this change would be zero.

5.2.5. Specific income processes. In general the difference in expectations which occur on the right of (28) depends on all the information which becomes available to the consumer between time t and $t + 1$. One special case which we now consider is when the only source of new information is the realisation of the labour income y_{t+1} . What revision this new information causes in the expectation of future labour income depends on how labour income in different periods are related.

As an example consider the labour income process given by the following stochastic difference equation,

$$y_{t+1} - \mu = \rho(y_t - \mu) + \epsilon_{t+1} \quad (29)$$

where ϵ_t is a white-noise process, μ and ρ are constants and $0 < \rho < 1$. This is a special case of what is known as a first-order autoregressive process (sometimes denoted as a AR(1) process). The coefficient ρ measures how persistent the deviations in y from its long-run average μ are.

Writing (29) for the period $t + 2$ we have,

$$y_{t+2} - \mu = \rho(y_{t+1} - \mu) + \epsilon_{t+2}$$

Substituting (29),

$$y_{t+2} - \mu = \rho^2(y_t - \mu) + \rho\epsilon_{t+1} + \epsilon_{t+2}$$

Carrying out successive substitutions like this, we find for any $i > t$

$$y_i - \mu = \rho^{i-t}(y_t - \mu) + \sum_{j=t+1}^i \rho^{i-j}\epsilon_j$$

Taking expectations conditional on the information at time t ,

$$E_t(y_i - \mu) = \rho^{i-t}(y_t - \mu) + \sum_{j=t+1}^i \rho^{i-j} E_t \epsilon_j$$

Here we have used the fact that y_t is known at time t . We further note that since ϵ_t is IID, ϵ_j is independent of all information at time t when $j > t$ and we can replace $E_t \epsilon_j$ by $E \epsilon_j$ which is 0 by the definition of white noise. Hence we conclude,

$$E_t(y_i - \mu) = \rho^{i-t}(y_t - \mu) \quad \text{for } i \geq t \quad (30)$$

(We have established this above for $i > t$ and it is trivially true for $i = t$.)

Using $t + 1$ in the place of t ,

$$E_{t+1}(y_i - \mu) = \rho^{i-t-1}(y_{t+1} - \mu) \quad \text{for } i \geq t + 1 \quad (31)$$

For $i \geq t + 1$ both (30) and (31) hold. Subtracting the former from the latter we have,

$$E_{t+1} y_i - E_t y_i = \rho^{i-t-1}(y_{t+1} - \mu) - \rho^{i-t}(y_t - \mu)$$

Using (29),

$$\begin{aligned} &= \rho^{i-t-1}[\rho(y_t - \mu) + \epsilon_{t+1}] - \rho^{i-t}(y_t - \mu) \\ &= \rho^{i-t-1} \epsilon_{t+1} \end{aligned}$$

Substituting this in (28) we get,

$$c_{t+1} - c_t = \left[\frac{\sum_{i=t+1}^T (\delta \rho)^{i-t-1}}{\sum_{i=t+1}^T \delta^{i-t-1}} \right] \epsilon_{t+1}$$

So we see that for a given innovation in consumption, ϵ_{t+1} , the increment in consumption is higher the higher is the degree of persistence ρ in the income process.

In empirical applications of the model we can estimate ρ (or its analogues for more complex income processes) from data on consumers' incomes and then check if changes in consumption satisfy the formula above. This yields a sharper test of our theory compared to just checking if consumption is a martingale.

AMBEDKAR UNIVERSITY DELHI

E-mail address: jyotirmoy@jyotirmoy.net