Macroeconomics

Jyotirmoy Bhattacharya

Ambedkar University Delhi
E-mail address: jyotirmoy@jyotirmoy.net

Contents

Consumption: Certainty

1. Two-period case

1.1. Budget constraint. Consider a consumer who lives for two periods, has an endowment of y_1 and y_2 units of goods in the two periods respectively and can borrow and lend any amount that they like at the real rate of interest r.

Suppose the consumer consumes c_1 in the first period. Then she will have to take a loan of $c_1 - y_1$ to finance her consumption. (This number can be negative, in which case the consumer is lending rather than borrowing.) In the next period the consumer will therefore have to make loan repayments of $(1+r)(c_1 - y_1)$. Assume that the consumer does not want to make any bequests and cannot die with any outstanding loans, consumption in the second period must be,

$$c_2 = y_2 - (1+r)(c_1 - y_1)$$

Simplifying and rearranging we have

$$c_1 + \frac{c_2}{1+r} = y_1 + \frac{y_2}{1+r} \tag{1}$$

This is the budget constraint faced by the consumer. We can interpret this to mean that the present value of the consumer's consumption stream must equal the present value of their incomes.

1.2. Utility maximization. Suppose the consumer maximises a quasiconcave utility function $U(c_1, c_2)$ subject to this budget constraint. Then the consumer's first-order conditions are

$$U_1(c_1, c_2) = \lambda \tag{2}$$
$$U_2(c_1, c_2) = \lambda/(1+r) \tag{3}$$

where λ is the Lagrange multiplier corresponding to the budget constraint and $U_i(c_1, c_2)$ denotes the partial derivative $\partial U/\partial c_i$. We have explicitly shown the dependence of the partial derivatives on the value of consumption in both periods. These first-order conditions along with the budget constraint (1) together determines the value of c_1, c_2 and λ.

1.3. Comparative statics. Assuming that consumption in both periods is a normal good, an increase in either y_1 or y_2 increases both c_1 and c_2.

The effects of a change in r are ambiguous. An increase in r makes consumption in period 2 relatively cheap compared to consumption in period 1. Therefore the substitution effect causes c_1 to decrease and c_2 to increase. It is traditional to decompose the income effect into two parts. First, an increase in r reduces the present value of the consumer's endowments and hence decreases his real income. Second, an increase in r, by making the consumption in period 2 cheaper increases his real income.[1] The sign of the resultant of these two effects on consumption depends on whether the consumer is a net lender in period 1 and a net borrower in period 2 or vice-versa. In case the consumer is a net lender in period 1 and a net borrower in period 2 the net income effect is positive. Assuming the consumption in both periods in a normal good, this means that the substitution effect and the income effect act in opposite directions on c_1 in this case leading to an ambiguous effect.

2. Many periods

Assume that rather than just living for two periods the consumer lives for $T + 1$ periods. Further assume that the real rate of interest takes a constant value r over the consumer's lifetime. For convenience we define $\delta = 1/(1+r)$. It is also convenient to start time from period 0 rather than period 1.

2.1. Budget constraint. Arguing as before, the consumer's budget constraint is

$$\sum_{i=0}^{T} \delta^i c_i = \sum_{i=0}^{T} \delta^i y_i \tag{4}$$

2.2. Utility function. We could proceed as before by assuming a utility function $U(c_0, \ldots, c_T)$ and deriving the first order conditions. However, because the marginal utility in each period depends on consumption in all periods it is hard to draw any sharp conclusions at this level of generality. Therefore we need to impose some restrictions on the form of the utility functions.

Suppose, for example we assume that the utility function is additively separable, i.e.

$$U(c_0, \ldots, c_T) = v_0(c_0) + v_1(c_1) + \cdots + v_T(c_T) \tag{5}$$

---

[1] For more about the Slutsky equation in the case of a consumer with fixed endowments of goods see section 9.1 in Varian's *Microeconomic Analysis*, 3rd ed.

Then the first-order conditions take the form

$$v_i'(c_i) = \delta^i \lambda \qquad i = 0, \ldots, T \tag{6}$$

where, as before, λ is the Lagrange multiplier corresponding to the budget constraint.

Sometimes we want to restrict the consumers preferences even further, by assuming that the different v_i differ from each other by only a geometric discounting factor.

$$U(c_0, \ldots, c_T) = \sum_{i=0}^{T} \beta^i u(c_i) \tag{7}$$

where β is a constant, referred to as the subjective rate of discount, such that $0 < \beta < 1$.

In this case the first-order conditions take the particularly simple form

$$u'(c_i) = \left(\frac{\delta}{\beta} \right)^i \lambda \qquad i = 0, \ldots, T \tag{8}$$

In case $\delta = \beta$, this implies that $u'(c_i)$ is the same for all i, which, assuming that $u'(\cdot)$ is a strictly decreasing function, means that c_i is constant for all i. The present period's income does not influence the present period's consumption at all. Consumption is determined solely by lifetime resources as given by (4).

The case $\delta \neq \beta$ is also instructive. Suppose $\delta > \beta$. In this case it follows from (8) that consumption decreases over time. Formally, this is because if $\delta > \beta$ then by (8) $u'(c_i)$ increases over time, and since $u'(c)$ is a decreasing function of consumption, this implies that c decreases over time.

The economic logic behind this result is that δ is the number of units of consumption we have to give up at present in order to purchase one more unit of consumption next period, whereas β is the number of units of marginal utility we are willing to give up at present in order to have one more unit of marginal utility in the next period. Suppose we start with the same consumption c in this period and the next. If we reduce consumption in the next period by a small amount Δc then at the prevailing market prices we can increase present consumption by $\delta \Delta c$. The increase in utility from the increase in present consumption is approximately $u'(c)(\delta \Delta c)$.[2] The decrease in utility from the reduction in next period's consumption is approximately $\beta u'(c)(\Delta c)$. The net change in utility would be $(\delta - \beta) u'(c)(\Delta c)$ which is positive when $\delta > \beta$.

---

[2]We are using Taylor's theorem: $u(c + \delta \Delta c) - u(c) \approx u'(c)(\delta \Delta c)$

Thus it is beneficial to increase present consumption and reduce future consumption if we are starting from a position of equality. Indeed, it will be optimal to increase consumption in the present period (say period i) and decrease consumption in the next period (period $i+1$) till the following equality between the MRS and the price ratio is satisfied,

$$\frac{u'(c_{i+1})}{u'(c_i)} = \frac{\delta}{\beta}$$

If δ/β is close to 1 then c_{i+1} is close to c_i and we can use Taylor's Theorem from calculus to the above equation to the above equation to get a useful approximation.

$$\frac{u'(c_{i+1})}{u'(c_i)} = \frac{\delta}{\beta}$$

$$\frac{u'(c_{i+1}) - u'(c_i)}{u'(c_i)} = \frac{\delta}{\beta} - 1$$

Applying Taylor's Theorem

$$\frac{u''(c_i)(c_{i+1} - c_i)}{u'(c_i)} \approx \frac{\delta}{\beta} - 1$$

Defining $\Delta c = c_{i+1} - c_i$, and dropping the subscript i,

$$\left(\frac{u''(c)c}{u'(c)}\right)\left(\frac{\Delta c}{c}\right) \approx \frac{\delta}{\beta} - 1$$

The quantity $\sigma = -u'(c)/cu''(c)$ is known as the *intertemporal elasticity of substitution* and captures the sensitivity of marginal utility of changes in consumption. It is positive since marginal utility decreases with consumption.

$$\left(\frac{\Delta c}{c}\right) \approx \sigma\left(1 - \frac{\delta}{\beta}\right)$$

The formula confirms our earlier reasoning that consumption decreases over time if $\delta > \beta$. Moreover, it shows that the sensitivity of the growth of consumption on the rate of return depends on the intertemporal elasticity of substitution. This is because the intertemporal elasticity of substitution is the reciprocal of the elasticity of marginal utility with respect to the level of consumption. The more elastic is marginal utility to consumption, the smaller is the deviation in consumption from a constant path that is required the equate the ratio of marginal utilities in consecutive time periods to δ/β.

2.3. Exogenous variables. It is possible to unify (5) and (7) by writing

$$v_i(c_i) = \beta^i u(c_i, \xi_i)$$

where ξ_i is an exogenous variable such a the consumer's age or the number of members in the household. In this case the first-order conditions become

$$u'(c_i, \xi_i) = \left(\frac{\delta}{\beta}\right)^i \lambda \qquad i = 0, \dots, T$$

Knowing how ξ affects the marginal utility would now let us make some predictions regarding the path of consumption.

2.4. Comparative statics. Assuming that consumption in every period is a normal good, an increase in y_i increases every c_i.

The effect of an increase in r, or equivalently, a decrease in δ remains ambiguous because of the same income and substitution effects as discussed earlier. But for the utility function given by (7), we can say a little more. From (8) we can see that a decrease in δ means that the *growth rate* of consumption speeds up. Remember that even in this case we do not have any information regarding the *level* of consumption in any period since the level would depend on λ which in turn depends on δ.

The Envelope Theorem

1. Parametrised optimisation problems

Let's think of unconstrained problems first. Every optimisation problem has an objective function. It is the function that we are trying to maximise or minimise (henceforth maximise). Some of the variables entering the objective function are *choice variables*, variables whose values we are free to choose in order to maximise the objective function. But all the variables entering into the objective function need not be choice variables. The value of the objective function may also depend on the value of other variables which we are not free to choose. We call these the *parameters* of the optimisation problem.

EXAMPLE 2.1. Consider the short-term profit maximising problem of a firm that produces according to the production function

$$y = f(L, K) = L^{1/2} K^{1/2}$$

In the short-run the capital stock of the firm is fixed at some value $\bar{K}$ and the firm can only choose the labour input L. If the firms buys labour and capital in perfectly competitive labour market at prices w and r respectively and sells its output in a perfectly competitive market at the price p then its profits are:

$$\pi(L, \bar{K}) = py - wL - r\bar{K} = pL^{1/2}\bar{K}^{1/2} - wL - r\bar{K}$$

For the short-run profit maximising problem $\pi(L, \bar{K})$ is the objective function, with L as a choice variable and $\bar{K}$ as a parameter.[1]

---

[1] In fact p, r and w are also parameters in the profit function. But we shall ignore this fact for now since we will not be looking at the effects of changes in these variables.

Denoting the optimal amount of labour input by L^*, the first-order condition for profit maximisation is,

$$\frac{\partial \pi}{\partial L} = 0$$

$$\frac{1}{2}pL^{*-1/2}\bar{K}^{1/2} - w = 0$$

$$L^* = \bar{K}(p/2w)^2$$

You should check that $\pi(L, \bar{K})$ is a concave function of L and therefore the first-order condition is sufficient to give us a global maximum. The profit earned by the firm at the optimal point is,

$$\pi^* = \pi(L^*, \bar{K})$$
$$= p[\bar{K}^{1/2}(p/2w)]\bar{K}^{1/2} - w[\bar{K}(p/2w)^2] - r\bar{K}$$
$$= \bar{K}(p^2/2w) - \bar{K}(p^2/4w) - r\bar{K}$$
$$= \bar{K}(p^2/4w) - r\bar{K}$$

We see that both the amount of labour input chosen by the firm and the maximum profit it earns are functions of the value of the parameter $\bar{K}$. The function mapping the parameter values to the maximum (or minimum) value of the objective function is called the *value function*. In this case, denoting the value function by $V(\cdot)$ we have,

$$V(\bar{K}) = \pi^* = \bar{K}(p^2/4w) - r\bar{K}$$

$\square$

2. The envelope theorem

How does the optimal value change when we change the parameters? In our example since we have an explicit formula for the value function we can calculate its value directly

$$V'(\bar{K}) = (p^2/4w) - r\bar{K}$$

Even when we do not have an explicit formula for the value function, there is an interesting relationship between the partial derivatives of the objective function and the partial derivatives of the value function.

Consider the general problem of maximising the objective function

$$\phi(x_1, \ldots, x_n; c_1, \cdots, c_m)$$

where the x_i are choice variables and c_i are parameters.

The first order conditions for the problem are,

$$\frac{\partial \phi}{\partial x_i}(x_1, \ldots, x_n; c_1, \ldots, c_m) = 0 \qquad i = 1, \cdots, n \qquad (9)$$

Just as in the example, the optimal values of the choice variables, denoted by x_i^*, will be functions of the parameters $c_1, \ldots, c_m$. The value function will be given by

$$V(c_1, \ldots, c_m) = \phi(x_1^*, \ldots, x_n^*; c_1, \ldots, c_m)$$

Suppose we want to calculate the partial derivative of the value function with respect to one of the parameters, say c_j. In doing so we have to take into account the fact that the optimal value of each of the choice variables would also be a function of c_i. So we use the chain rule,

$$\frac{\partial V}{\partial c_j} = \frac{\partial \phi}{\partial x_1} \frac{\partial x_1^*}{\partial c_j} + \cdots + \frac{\partial \phi}{\partial x_n} \frac{\partial x_n^*}{\partial c_j} + \frac{\partial \phi}{\partial c_j}$$

However, from (9), we know that $\partial \phi / \partial x_i$ is 0 for all i when the partial derivatives are evaluated at the optimal values. So we have,

$$\frac{\partial V}{\partial c_j} = \frac{\partial \phi}{\partial c_j} \tag{10}$$

This remarkably is the same result that we would have got if we had treated each of the x_i^* as a constant. But that would not have been justified since the choice variables do vary when parameters are varied. That is, $\partial x_i^* / \partial c_j$ is generally not zero. It is just that when we are starting from an optimal point then the marginal impact on this variation on the objective function (i.e., $\partial \phi / \partial x_i$ is zero and therefore we can ignore the changes in the choice variables.

Equation (10) is known as the "Envelope Theorem".

3. Geometric Interpretation

Figure 1 illustrates the envelope theorem in the case of Example 2.1. Each of the coloured curves shows the level of profit for a given level of L and for different values of K. Let's call them "profit curves".[2] We have drawn only three of these curves but you should imagine there to be one curve for each possible value of L. Now, since our purpose is to maximise profit for a given value of K, we move along a vertical line for our particular value of K and choose that L whose profit curve is the highest at that value of K.

Thus, for example, at $K = 4.0$ we would choose $L = 4.0$ whereas at $K = 10.0$ we would choose $L = 2.5$.

The value of the highest profit curve for a given K gives us the highest profit we can obtain when K takes on that value. But that

---

[2]This is not standard terminology and you must remember that these curves are not graphs of the full profit function since we are holding L constant on each of them.

FIGURE 1. The Envelope Theorem

is precisely the definition of the value function. Therefore the graph
of the value function touches the highest of the profit curves at each
K. Or, in other words, the graph of the value function (the black line
in the figure) must be the upper envelope of the graphs of the profit
functions for given values of L.

Since the value function is the upper envelope of the profit curves,
no profit curve can ever cross it. But at each value of K one of the
profit curves, corresponding to the optimal L, touches it. The only
way two graphs can touch without crossing is if they are tangent to
each other. The slope of the graph of the value function is $\partial V/\partial K$
whereas the slope of the profit curves is $\partial \pi/\partial K$. Tangency of the two
graphs implies that these slopes should be equal, which is precisely
what our the envelope theorem in eq. (10) also says when applied to
this example.

Now you know what the envelope theorem is called by that name.

4. Constrained Optimisation

So far we have discussed unconstrained problems. There is also a
version of the envelope theorem for constrained optimisation problems.
Suppose our problem is to maximise

$$\phi(x_1, \ldots, x_n; c_1, \ldots, c_m)$$

subject to the constraint

$$h(x_1, \ldots, x_n; c, \ldots, c_m) = 0 \tag{11}$$

Here we have allowed both the objective function and the constraint to depend on a set of parameters.

The first-order condition for this problem is

$$\frac{\partial \phi}{\partial x_i} = \lambda \frac{\partial h}{\partial x_i} \qquad i = 1, \ldots n \tag{12}$$

where λ is a Lagrange multiplier.

As before, if the problem has a solution the optimal values of the choice variables, the x_i^*, will be functions of the parameters of the problem. Also as before, we can define the value function as

$$V(c_1, \ldots, c_m) = \phi(x_1^*, \ldots, x_n^*; c_1, \ldots, c_m)$$

Differentiating the value function with respect to c_j gives us,

$$\frac{\partial V}{\partial c_j} = \frac{\partial \phi}{\partial x_1} \frac{\partial x_1^*}{\partial c_j} + \cdots + \frac{\partial \phi}{\partial x_n} \frac{\partial x_n^*}{\partial c_j} + \frac{\partial \phi}{\partial c_j} \tag{13}$$

To simplify this we need to digress a bit. The optimal values of the choice variables must satisfy the constraint (11) for all values of the parameters, so we have

$$h(x_1^*, \ldots, x_n^*; c, \ldots, c_m) = 0.$$

Differentiating this with respect to c_j we get

$$\frac{\partial h}{\partial x_1} \frac{\partial x_1^*}{\partial c_j} + \cdots + \frac{\partial h}{\partial x_n} \frac{\partial x_n^*}{\partial c_j} + \frac{\partial h}{\partial c_j} = 0$$

Substituting the first-order conditions (12) we have,

$$\frac{1}{\lambda} \frac{\partial \phi}{\partial x_1} \frac{\partial x_1^*}{\partial c_j} + \cdots + \frac{1}{\lambda} \frac{\partial \phi}{\partial x_n} \frac{\partial x_n^*}{\partial c_j} + \frac{\partial h}{\partial c_j} = 0$$

or,

$$\frac{\partial \phi}{\partial x_1} \frac{\partial x_1^*}{\partial c_j} + \cdots + \frac{\partial \phi}{\partial x_n} \frac{\partial x_n^*}{\partial c_j} = -\lambda \frac{\partial h}{\partial c_j}$$

Now this can be substituted in (13) to give us

$$\frac{\partial V}{\partial c_j} = -\lambda \frac{\partial h}{\partial c_j} + \frac{\partial \phi}{\partial c_j} \tag{14}$$

Equation (14) is the envelope theorem for the constrained case. It is similar to the unconstrained envelope theorem in that the change in the choice variables as a result of the change in the parameters drops out of the calculation. It differs in that the change in the value function as a result of a change in a parameter depends not just on the direct change

in the objective function $(\partial \phi / \partial c_j)$ but also the change in constraint set $(\partial h / \partial c_j)$. The Lagrange multiplier λ can be interpreted as a sensitivity factor, indicating the extent to which a given change in the constraint set translates into a change in the value function.

EXAMPLE 2.2. Consider the problem of maximising the utility function $U(x_1, x_2)$ subject to the budget constraint $p_1 x_1 + p_2 x_2 = M$. Treating x_1 and x_2 as choice variables and p_1, p_2 and M as parameters, we have the objective function

$$\phi(x_1, x_2; p_1, p_2, M) = U(x_1, x_2)$$

and the constraint function

$$h(x_1, x_2; p_1, p_2, M) = p_1 x_1 + p_2 x_2 - M$$

In this case the value function $V(p_1, p_2, M)$ is important enough to be given a name. It is called the *indirect utility function.*

If the value of the Lagrange multiplier at a the optimal bundle is λ, then the envelope theorem (14) tells us that,

$$\frac{\partial V}{\partial M} = -\lambda \frac{\partial h}{\partial M} + \frac{\partial \phi}{\partial M}$$
$$= -\lambda \cdot -1 + 0$$
$$= \lambda$$

This gives us an economic interpretation of the Lagrange multiplier. It measures the amount by which the maximum attainable utility increases per unit increase in income. In looser phrasing, it is the "marginal utility of income".

$\square$

CHAPTER 3

Dynamic programming

1. The setup

In a dynamic optimisation problem, our goal is to find a *path* of the choice variable which maximises the value of an objective function defined over the entire path of the choice variable. Often, there are constraints on what paths can be chosen. For example, in the consumption-saving problem we choose a path of consumption which maximises the lifetime utility function subject to a budget constraint.

The dynamic programming approach to solving dynamic optimisation problems turns this single large optimization problem into a sequence of simple optimization problems. At each point of time we try to find the best action at that particular point of time. But since this is a dynamic problem after all, this search for the best actions at a particular point of time has to be done with an eye on both the past and the future. Past events and actions[1] determine what choices can be made now. By the same token, the action that we take now will change the options available to us the in future. The value of the objective function that will be achieved will in general depend on the entire path of past, present and future actions and not just the action in any period in isolation.

In the dynamic programming framework this linkage between the past and the future is captured by the notion of the *state*. Intuitively, we can think of the state at any given point of time as a description of all the *relevant* information about the actions and events that have happened until that point. The state should contain all the information that is required from the decision-maker's history to determine the set of available actions at future points of time and to evaluate the contribution made by future actions to the objective function.

The notion that knowing the state at a point of time is enough to know what actions are possible in the future is captured by the following definitions:

---

[1]We want to make a distinction between *events* which are outside of our control and *actions* which are things we choose. This distinction becomes important when we are dealing with uncertainty.

Set of states (S_t)**:** this is the set of possible states the decision-maker can be in time t. There is one such set for each time period t. The elements of these sets, i.e. the possible states, are assumed to be vectors with real-number elements. Elements of this set are denoted by s_t.

Set of actions (A_t)**:** this is the set of possible actions that can be taken at time t. Elements of this set are denoted by a_t. As we shall see next, all possible actions cannot be taken at all possible states.

Constraint correspondence $f_t(s_t) \subset A_t$**:** this tells us the subset of actions that are available in a particular state. This is not a function but a correspondence (i.e. a set-valued function) since for each element of S_t it gives us a subset and not just a single element of A_t.

Transition function $\Gamma_t(s_t, a_t) \in S_{t+1}$**:** This tells us our state in period $t+1$ if we take the action a_t in state s_t in period t.

With these definitions in hand we can define the set of *feasible plans* when starting with s_t at time t, denoted by $\Phi_t(s_t)$, as the set of sequences of actions $(a_t, \dots, a_T)$ such that

$$a_i \in f_i(s_i) \quad \text{for } i = t, \dots, T$$

and

$$s_{i+1} = \Gamma_i(s_i, a_i) \quad \text{for } i = t, \dots, T-1$$

The first condition says that the action taken on each date is a feasible action given the state. The second condition says that the state at each date is derived from the state and action taken in the previous date, with the state at time t as given.

The set of feasible plans tells us about the constraints faced in our optimisation problem. What about the objective function? We assume that the objective function can be written in an additively separable form

$$U_t(a_t, s_t, \dots, a_T, s_T) = \sum_{i=t}^{T} v_i(a_i, s_i)$$

where $v_t(a_t, s_t)$ is the *per-period payoff function* that gives the contribution of action a_t in state s_t at time t to the overall objective. Being able to write the objective function in an additively separable form is essential for us to be able to use dynamic programming.[2]

_____

[2]We are cheating a bit here. The assumption of additive separability can be relaxed to what is called 'recursiveness' while still allowing the use of dynamic programming.

In writing the above objective function we have also assumed that there is a finite time period T at which our optimisation problem comes to an end. This assumption of what is known as a finite horizon is made just to simplify the mathematics. Dynamic programming problems with an infinite horizon are routinely used in economic modelling.

The solution to the dynamic programming problem is expressed in terms of two functions:

Policy function $g_t(s_t) \in f_t(s_t)$**:** The policy function tells us the best action to take in each possible state at time t among all the available actions. In general it is possible that there be two equally good actions at a particular state, in which case the policy function would have to be replaced by the policy correspondence.

Value function $V_t(s_t) \in \mathbb{R}$**:** The value function denotes the maximum attainable value of the objective function when starting at time t from state s_t. That is,

$$V_t(s_t) = \max_{(a_t,\ldots,a_T)\in\Phi_t(s_t)} U_t(a_t, s_t, \ldots, a_T, s_T)$$

$$= \max_{(a_t,\ldots,a_T)\in\Phi_t(s_t)} [v_t(a_t, s_t) + \cdots + v_T(a_T, s_T)]$$

In applications of dynamic programming we generally want to know the optimal path starting at a specific point of time (taken to be $t = 0$ here) and from a particular state at that point of time (say $\bar{s}_0$). But we have defined the policy and value functions for all points of time and for each possible state at each of the time periods. Thus it would seem that we have multiplied our work manyfold beyond what is necessary for our original problem. But as we shall see below, being willing to contemplate the policy and value functions for all possible time periods and states often actually simplifies the task of solving the original problem.

2. Bellman's Principle of Optimality

Suppose I am starting at some time $t < T$ from some particular state s_t and trying to find the best actions from time t to T, where 'best' means the choices of actions and consequent states which maximise

$$U_t(a_t, s_t, \ldots, a_T, s_T) = \sum_{i=t}^{T} v_i(a_i, s_i).$$

Because of the additive nature of the lifetime utility function we can rewrite the above equation as

$$U_t(a_t, s_t, \ldots, a_T, s_T) = v_t(a_t, s_t) + U_{t+1}(a_{t+1}, s_{t+1}, \ldots, a_T, s_T)$$

If we divide the plan (path of actions and corresponding states) from time t to time T into a "head" consisting of the action in period t and a "tail" consisting of actions in period $t+1$ to T then the above equation says that the lifetime utility of the plan starting at period t is the sum of the per-period payoff at time t (the value of the "head") and the lifetime utility of the remaining part of the plan from period $t+1$ onwards (the value of the "tail").

How do we find the plan which maximises U_t? Suppose we choose the action $\hat{a}_t$ in period t. This will lead us to the state $\hat{s}_{t+1} = \Gamma_t(s_t, \hat{a}_t)$ in the next period. Now we have to pick a plan from period $t+1$ onward. Now $U_t = v_t(\hat{a}_t, s_t) + U_{t+1}$ and $v_t(\hat{a}_t, s_t)$ is already fixed by our choice of action $\hat{a}_t$ in period t. Therefore in choosing our plan from period $t+1$ onward the best we can do is to pick a plan that maximises U_{t+1}. This optimal plan for the "tail" yields the value of U_{t+1} equal to $V_{t+1}(\hat{s}_{t+1})$. Thus we can evaluate each choice of action $\hat{a}_t$ in the "head" by looking at

$$\tilde{U}_t = v_t(\hat{a}_t, s_t) + V_{t+1}(\hat{s}_{t+1}), \quad \text{where } \hat{s}_{t+1} = \Gamma(s_t, \hat{a}_t)$$

We have put a tilde over U_t to remind ourselves that now we are not considering arbitrary plans starting at t but only plans where the "tail" component is optimal given the state $\hat{s}_{t+1}$ at which we find ourselves in the beginning of period $t+1$.

The optimal plan from period t involves choosing $\hat{a}_t$ which maximises the expression above. Since the value function gives the value of the objective function U_t for the optimal plan, it is therefore the case that,

$$V_t(s_t) = \max_{\hat{a}_t \in f_t(s_t)} [v_t(\hat{a}_t, s_t) + V_{t+1}(\Gamma(s_t, \hat{a}_t))], \quad \text{for } t < T \qquad (15)$$

Equation (15) above which relates the value function at consecutive time periods is known as *Bellman's Equation*. The argument above, which shows that the value function must satisfy Bellman's equation is known as *Bellman's Principle of Optimality*.[3]

Intuitively Bellman's equation tells us that we can evaluate each present action by adding its contribution $v_t(\hat{a}_t, s_t)$ to the objective in the present period and the value $V_{t+1}(\hat{s}_{t+1})$ of the state in which it leaves us in the next period. Provided we know the function V_{t+1} for all possible states in the next period we can choose the best action in the current period by choosing $\hat{a}_t$ to maximise this sum. Thus we have

---

[3]To be complete, Bellman's principle of optimality also deals with the converse: that a function which satisfies Bellman's equation plus some other technical conditions must be the value function. This converse is not important in our current finite horizon setting.

turned the big optimisation problem of choosing an entire sequence of actions from time 0 to time T into a sequence of simple optimisation problems, one for each time period t, in each of which we choose a single action $\hat{a}_t$.

But there seems to be a chicken-and-egg problem: we cannot use Bellman's equation without knowing V_{t+1} for each t and how do we know V_{t+1} if we have not solved the optimisation problem already? Here our finite horizon assumption makes life particularly simple for us.

Since period T is the last period, our objective function in that period is

$$U_T(a_T, s_T) = v_T(a_T, s_T)$$

and the value function is simply given by

$$V_T(s_T) = \max_{a_T \in f_T(s_T)} v_T(a_T, s_T)$$

We can solve this maximisation problem and calculate V_T since $v_T(\cdot, \cdot)$ is a known function.

Now consider Bellman's equation for period $T - 1$:

$$V_{T-1}(s_{T-1}) = \max_{\hat{a}_{T-1} \in f_t(s_{T-1})} \left[v_{T-1}(\hat{a}_{T-1}, s_{T-1}) + V_T(\Gamma(s_{T-1}, \hat{a}_{T-1})) \right]$$

As we have calculated $V_T(\cdot)$ in the previous step, all the functions in the maximisation problem are known and we can solve the problem to calculate V_{T-1}. With this in hand we can solve Bellman's equation for period $T - 2$. We keep going backward one period at a time until we have calculated the value function for all periods until period 0. At each step the value of $\hat{a}_t$, as a function of s_t, which solves the maximisation problem gives us the policy function. So by the end of our process we also have the policy function for each time period.

Now if we are given a starting state $\bar{s}_0$ in period 0 we can use the calculated policy function for period 0 to find the best action a_0 in period 0. We know from the transition function that we will end up in state $s_1 = \Gamma(\bar{s}_0, a_0)$ in the next period. The policy function for period 1 tells us the best action a_1 to take in that period. We again use the transition function to tell us the next state $s_2 = \Gamma(s_1, a_1)$. And so on until we have traced out the optimal plan to time T. Our optimisation problem is solved!

The way we have calculated the value function backwards from a known final time period is sometimes called "backward induction".

3. Example: consumption-savings with log utility

Suppose that the consumer maximises

$$\sum_{i=0}^{T} \log(c_i)$$

subject to

$$\sum_{i=0}^{T} c_i/R^i = w_0$$

Can we solve this maximisation problem using dynamic programming? The action variable in this case must be c_i since it is the variable being chosen by the decision maker. The objective function is already in an additively separable form with a per-period payoff $\log(c_i)$. But what is the state?

Since the per-period payoff depends only on the action variable c_i we do not need any notion of state to evaluate payoffs. But the choice of a consumption in each period does affect future periods through the budget. The more we consume today, the less purchasing power we have to consume tomorrow. We can capture this by rearranging the budget slightly to read,

$$\sum_{i=1}^{T} c_i/R^i = w_0 - c_0$$

Multiplying throughout by R we have

$$\sum_{i=1}^{T} c_i/R^{(i-1)} = R(w_0 - c_0)$$

Which shows us that the path of consumption from time 1 onwards follows a budget constraint of the same form as the period 0 budget constraint provided we take

$$w_1 = R(w_0 - c_0)$$

This suggests to us that we can take the wealth at the beginning of period t as our state with the transition function,

$$w_{t+1} = R(w_t - c_t), \quad \text{for } t = 0, \ldots, T$$

and the constraint function

$$c_T = w_T$$

The constraint captures the fact that there can be no outstanding debt in the last period and a consumer who has monotonic preferences would

not leave any wealth unused in the last period. There is no constraint on consumption in periods other that T.[4].

You can check that we have formulated the problem right by eliminating $w_1, \ldots, w_T$ in the transition functions and constraints above to recover our original budget constraint.

Now we can use backward induction to calculate the value function and policy function for all time periods. Denoting the policy function by $g(\cdot)$ and the value function in period t by $V_t(\cdot)$, we have,

$$g_T(w) = w, \qquad V_T(w_T) = \log(g_T(w)) = \log(w) \qquad (16)$$

Now consider period $T - 1$. Bellman's principle of optimality tells us,

$$
\begin{aligned}
V_{T-1}(w) &= \max_c [\log(c) + V_T(R(w - c))] \\
&= \max_c [\log(c) + \log(R(w - c))] \qquad [\text{using } (16)]
\end{aligned}
\qquad (17)
$$

The first-order condition for this maximisation problem is:

$$\frac{1}{c} + \frac{-R}{R(w - c)} = 0$$
$$w - c = c$$
$$c = w/2$$

Since the objective function in (17) is concave in c (check this!), the first-order condition is sufficient and gives us our policy function:

$$g_{T-1}(w) = w/2$$

Substituting this into (17) we get the value function,

$$
\begin{aligned}
V_{T-1}(w) &= \log(g_{T-1}(w)) + V_T(R(w - g_{T-1}(w))) \\
&= \log(w/2) + \log(Rw/2) \\
&= \log(R) + 2\log(w/2)
\end{aligned}
\qquad (18)
$$

Now that we know V_{T-1} we could use the Bellman equation relating V_{t-2} to V_{t-1} to derive g_{T-2} and V_{t-2}. If you do this you will find,

$$g_{T-2}(w) = w/3, \qquad V_{T-2}(w) = (1 + 2)\log(R) + 3\log(w/3) \qquad (19)$$

We could continue like this to find $V_{T-3}, \ldots, V_0$. In general this is precisely what we do. In fact, in most applications of dynamic programming it is not possible to express the value function by a formula in the state variables and the best that we can do is to use a computer

---

[4]We could have imposed a non-negativity constraint but we leave it out for simplicity

to calculate the value function at a number of possible values of the state variable using Bellman's equation.

But our present problem is a particularly simple one. Looking at (18) and (19) suggests to us the guess,

$$V_{T-n}(w) = \frac{n(n+1)}{2}\log(R) + (n+1)\log\left(\frac{w}{n+1}\right) \qquad (20)$$

[Remember $1 + 2 + \cdots + n = n(n+1)/2$]

How do we check that our guess is right? We will use the principle of mathematical induction. By comparing to (18) we see that (20) is correct for $n = 1$. Suppose that the equation is true for $n = k$. What then would be $V_{T-(k+1)}$? We once again write down the Bellman equation

$$
\begin{aligned}
V_{T-(k+1)} &= \max_c[\log(c) + V_{T-k}(R(w-c))] \\
&= \max_c\left[\log(c) + \frac{k(k+1)}{2}\log(R) + (k+1)\log\left(\frac{R(w-c)}{k+1}\right)\right] \\
&\quad [\text{assuming } (20)]
\end{aligned}
$$

$$(21)$$

The first-order condition is:

$$
\begin{aligned}
\frac{1}{c} + (k+1)\left(\frac{k+1}{R(w-c)}\right)\left(\frac{-R}{k+1}\right) &= 0 \\
(k+1)\frac{1}{(w-c)} &= \frac{1}{c} \qquad (22) \\
c &= w/(k+2)
\end{aligned}
$$

Substituting this into (21) we have

$$
\begin{aligned}
V_{T-(k+1)} &= \log(c) + \frac{k(k+1)}{2}\log(R) + (k+1)\log\left(\frac{R(w-c)}{k+1}\right) \\
&\quad \text{substituting } (22), \\
&= \log\left(\frac{w}{k+2}\right) + \frac{k(k+1)}{2}\log(R) + (k+1)\log\left(\frac{Rw}{k+2}\right) \\
&= \log\left(\frac{w}{k+2}\right) + \frac{k(k+1)}{2}\log(R) + (k+1)\log(R) + (k+1)\log\left(\frac{w}{k+2}\right) \\
&= \frac{(k+2)(k+1)}{2}\log(R) + (k+2)\log\left(\frac{w}{k+2}\right)
\end{aligned}
$$

$$(23)$$

But this is the same as (20) for $n = k + 1$. We therefore conclude that if (20) is true for $n = k$ it is also true for $n = k + 1$. We have already checked that (20) is true for $n = 1$. Hence we conclude by the principle of mathematical induction that the value function for our dynamic programming problem is given by (20) for $n = 1, \ldots, T$.

Also, now that we have verified that (20) is indeed the value function of the problem, (22) gives the policy function, i.e.

$$g_{T-n}(w) = w/(n+1) \tag{24}$$

4. The Euler equation

As we discussed in the last section, for most dynamic programming problems it is not possible to compute the value and policy functions in terms of simple formulae. The best we can do is to calculate numerical values. But even if we cannot find an exact formula for the solution to our optimisation problem, it may still be possible to get some qualitative information about the problem by studying the consequences of the Bellman equation. That is the subject of this section.

Let's recall the Bellman equation,

$$V_t(w_t) = \max_{c_t}[u(c_t) + V_{t+1}(R(w_t - c_t))$$

The first order condition for this maximisation problem is:

$$u'(c_t) = RV'_{t+1}(R(w_t - c_t)) \tag{25}$$

By itself (25) does not seem very useful unless we know $V_{t+1}(\cdot)$ and can calculate its derivative. But there is a trick that we can use to eliminate this unknown derivative from (25).[5]

Let $c_t^*(w_t)$ be the optimal consumption in period t when period t wealth is w_t. From (25) we already know that,

$$u'[c_t^*(w_t)] = RV'_{t+1}[R(w_t - c_t^*(w_t))] = RV'_{t+1}(w_{t+1}) \tag{26}$$

But from the definition of the value function

$$V_t(w_t) = u[c_t^*(w_t)] + V_{t+1}[R(w_t - c_t^*(w_t))]$$

Differentiating with respect to w_t we have,

$$V'_t(w_t) = u'[c_t^*(w)]c_t^{*\prime}(w_t) + V'_{t+1}[R(w_t - c_t^*(w_t))][R(1 - c_t^{*\prime}(w_t))]$$
$$= c_t^{*\prime}(w_t)[u'(\cdot)) - RV'_{t+1}(\cdot)] + RV'_{t+1}[R(w_t - c_t^*(w_t))]$$

---

[5]The 'trick' is a particular case of a general result known as the envelope theorem. See section M.L of Mas-Colell, Whinston and Green or some mathematical methods book for more detail.

From (26) the first terms equals 0, so,

$$V'_t(w_t) = RV'_{t+1}(w_{t+1})$$

Using (26)

$$V'_t(w_t) = u'(c_t)$$

The equation above was derived for arbitrary t. So it is equally good for $t+1$, i.e.

$$V'_{t+1}(w_{t+1}) = u'(c_{t+1})$$

Substituting this in (26) we have,

$$\frac{u'(c_{t+1})}{u'(c_t)} = \frac{1}{R} = \delta \tag{27}$$

This condition is known as the Euler[6] equation for our dynamic programming problem. We can alternatively derive it by starting out with an optimal consumption plan, increasing consumption in period t by a small amount Δc and reducing consumption in period $t+1$ by $R\Delta c$ so that wealth at the end of the period $t+1$ is once again the same as what it would have been under the optimal plan. The first-order change in utility from this deviation is

$$\Delta u = u'(c_t)[\Delta c] - u'(c_{t+1})[R\Delta c]$$

Now for the original plan to have been optimal Δu must be 0 since if $\Delta u > 0$ the deviation considered above increases total utility whereas if $\Delta u < 0$ then the opposite of the deviation considered above increases total utility. But $\Delta u = 0$ implies

$$u'(c_t) - Ru'(c_{t+1}) = 0$$

which is again our Euler equation (27).

The Euler equation also follows from the first-order conditions (6) of the Lagrange-multiplier approach, showing that we have come full circle.

The Euler equation tells us how consumption should grow or decline. It does not tell us the level of the consumption. But we can characterise the entire consumption path if we keep track of the path of wealth implied by the path of consumption and impose, in addition to the Euler equation, the conditions

$$w_0 = \overline{w_0}$$

---

[6]Pronounced "oiler". Named after a eighteenth-century mathematician who was among the earliest to study dynamic optimisation problems.

which comes to us as a given data and

$$w_T = 0$$

which comes to us from our no bequest, no terminal borrowing, monotonic utility assumptions about the terminal period.

Probability

This chapter is not a substitute for studying a probability text such as [**CA03**] or [**Ros09**]. We only discuss some ideas important for macroeconomic applications which cannot be found easily in elementary books like the above.

The discussion below is in terms of discrete random variables. Random variables which are not discrete can be thought of as limits of discrete random variables. For example, we can think of GDP being measured in whole rupees or whole paise, and hence being a discrete random variable, rather than thinking of it as an infinitely divisible real number. In this way the intuition developed for discrete random variables can be carried over to general random variables too. But the details are tricky and dealing with them in a rigorous manner requires the use of techniques from an area of mathematics knows as "measure theory" which we cannot cover here. If you are interested in such things you can look at [**JP04**] for an introduction and [**Bil95**] for a classic but demanding treatment.

1. The Law of Iterated Expectations

1.1. Calculation. We can calculate the conditional expectation of a particular random variable for different sets of conditioning variables. For example, if Y_t is the GDP of a country in period t then

$$E[Y_{2020}|Y_{2010}]$$

is the conditional expectation of GDP in 2020 given information about GDP in 2010. This will be a function of the realised value of GDP in 2010.

Similarly,

$$E[Y_{2020}|Y_{2010}, Y_{2015}]$$

is the conditional expectation of GDP in 2020 but now given information about GDP both in 2010 and 2015. This will be a joint function of the realised values of GDP in 2010 and in 2015.

In examples like the above, where one of the sets of conditioning variables is a superset of another, we can derive a relationship between

the two conditional expectations known as the *law of iterated expectations*. To state the law in a general form we gather together the conditioning variables into a set called the *information set*. So the two conditional expectations above would be written as

$$E[Y_{2020}|I_1], \quad I_1 = \{Y_{2010}\}$$

and

$$E[Y_{2020}|I_2], \quad I_2 = \{Y_{2010}, Y_{2015}\}$$

respectively where I_1 and I_2 are the information sets.

The law of iterated expectations says that for any random variable Y

$$E[E[Y|I_2]|I_1] = E[Y|I_1], \quad \text{if } I_1 \subset I_2 \tag{28}$$

The inner expectation $E[Y|I_2]$ is a function of the realised values of the random variables in the information set I_2. But suppose we don't know the values of all the variables in I_2 but only those of a subset I_1. Then the left-hand side of (28) is asking us to carry out a two-step process. First we calculate the conditional expectation $E[Y|I_2]$ for all possible realisations of variables in I_2. Then we average out these conditional expectations using the probability of the realisation of the different values of I_2 conditional on variables in I_1. This is $E[E[Y|I_2]|I_1]$. The law of iterated expectations tells us that the result of this two-step process is the same as that of directly calculating the expectation of Y conditional on I_1.

EXAMPLE 4.1. Suppose I toss an unbiased coin thrice. If the coin comes up heads I win one rupee and if the coin comes up tails I lose one rupee. Let Y_1, Y_2 and Y_3 be my cumulative net winnings after the first, second and third tosses respectively.

The probability mass function (pmf) of Y_3 conditional on Y_1 and Y_2 is

$$f(y_3|y_1, y_2) = \begin{cases} 1/2 & y_3 = 1, y_1 = 1, y_2 = 0 \\ 1/2 & y_3 = -1, y_1 = 1, y_2 = 0 \\ 1/2 & y_3 = 3, y_1 = 1, y_2 = 2 \\ 1/2 & y_3 = 1, y_1 = 1, y_2 = 2 \\ 1/2 & y_3 = 1, y_1 = -1, y_2 = 0 \\ 1/2 & y_3 = -1, y_1 = -1, y_2 = 0 \\ 1/2 & y_3 = -1, y_1 = -1, y_2 = -2 \\ 1/2 & y_3 = -3, y_1 = -1, y_2 = -2 \\ 0 & \text{elsewhere} \end{cases}$$

The expectation of Y_3 conditional on Y_1 and Y_2 is therefore given by,

$$E[Y_3|Y_1,Y_2] = \begin{cases} 1/2 \cdot 1 + 1/2 \cdot -1 = 0 & y_1 = 1, y_2 = 0 \\ 1/2 \cdot 3 + 1/2 \cdot 1 = 2 & y_1 = 1, y_2 = 2 \\ 1/2 \cdot 1 + 1/2 \cdot -1 = 0 & y_1 = -1, y_2 = 0 \\ 1/2 \cdot -1 + 1/2 \cdot -3 = -2 & y_1 = -1, y_2 = -2 \end{cases}$$

Now suppose that instead of observing both Y_1 and Y_2 we observe only Y_1. To carry out the two-step process required by the law of iterated expectations we need the pmf of Y_2 conditional on Y_1.

$$f(y_2|y_1) = \begin{cases} 1/2 & y_2 = 0, y_1 = 1 \\ 1/2 & y_2 = 2, y_1 = 1 \\ 1/2 & y_2 = 0, y_1 = -1 \\ 1/2 & y_2 = -2, y_1 = -1 \\ 0 & \text{elsewhere} \end{cases}$$

So we have

$$E[E[Y_3|Y_1,Y_2]|Y_1] = \sum_{y_2} E[Y_3|Y_1 = y_1, Y_2 = y_2] f(y_2|y_1) \tag{29}$$

$$= \begin{cases} 1/2 \cdot 0 + 1/2 \cdot 2 = 1 & y_1 = 1 \\ 1/2 \cdot 0 + 1/2 \cdot -2 = -1 & y_1 = -1 \end{cases} \tag{30}$$

On the other hand if we had wanted to directly calculate $E[Y_3|Y_1]$ we would have needed the pmf of Y_3 conditional on Y_1 which is

$$f(y_3|y_1) = \begin{cases} 1/4 & y_3 = 3, y_1 = 1 \\ 1/2 & y_3 = 1, y_1 = 1 \\ 1/4 & y_3 = -1, y_1 = 1 \\ 1/4 & y_3 = 1, y_1 = -1 \\ 1/2 & y_3 = -1, y_1 = -1 \\ 1/4 & y_3 = -3, y_1 = -1 \\ 0 & \text{elsewhere} \end{cases}$$

So the conditional expectation is

$$E[Y_3|Y_1] = \begin{cases} 1/4 \cdot 3 + 1/2 \cdot 1 + 1/4 \cdot -1 = 1 & y_1 = 1 \\ 1/4 \cdot 1 + 1/2 \cdot -1 + 1/4 \cdot -3 = -1 & y_1 = -1 \end{cases} \tag{31}$$

Comparing (29) and (31) we verify that the law of iterated expectations holds in this case. □

1.2. Intuition. One can think of an expectation as a weighted average with the probabilities of outcomes acting as the weights. In the case of an unconditional expectation we take the average over the entire sample space. In the case of a conditional expectation we break up the sample space into chunks, each chunk consisting of one possible set of realisations of the variables in the information set, and then average separately over each chunk.

Thus, in Example 4.1 the sample space is

$$\Omega = \{HHH, HHT, HTH, HTT, THH, THT, TTH, TTT\}.$$

If we condition on the information set $I_1 = \{Y_1\}$ then the sample space is broken into the following chunks

Chunk	Realisation of I_1
$C_1 = \{HHH, HHT, HTH, HTT\}$	$Y_1 = 1$
$C_2 = \{THH, THT, TTH, TTT\}$	$Y_1 = -1$

TABLE 1. Decomposition of the sample space by the information set I_1

On the other hand, if the information set is $I_2 = \{Y_1, Y_2\}$ then the chunks are

Chunk	Realisation of I_2
$C_1' = \{HHH, HHT\}$	$Y_1 = 1, Y_2 = 2$
$C_2' = \{HTH, HTT\}$	$Y_1 = 1, Y_2 = 0$
$C_3' = \{THH, THT\}$	$Y_1 = -1, Y_2 = 0$
$C_4' = \{TTH, TTT\}$	$Y_1 = -1, Y_2 = -2$

TABLE 2. Decomposition of the sample space by the information set I_2

The important thing to note here is that each of the chunks created by I_1 is a union of some chunks created by I_2. For example, $C_1 = C_1' \cup C_2'$. This is a consequence of I_1 being a subset of I_2. So if two outcomes produce the same values for the variables in I_2 they must necessarily produce the same values for the variables in I_1.

In this setup $E[Y_3|I_2]$ averages Y_3 separately over each of the chunks $C_1', \ldots, C_4'$ whereas $E[Y_3|I_1]$ averages Y_3 separately over C_1 and C_2.

The law of iterated expectations (28) says, for example, that it does not matter in the calculation of $E[Y_3|Y_1 = 1]$ if we replace the actual values of Y_3 on HHH and HHT by their average over C_1' which is $E[Y_3|Y_1 = 1, Y_2 = 2]$ and the actual values of Y_3 on HTH and HTT by their average over C_2' which is $E[Y_3|Y_1 = 1, Y_2 = 0]$.

The intuitive reason is that all the outcomes in any C_i' are going to be in the same C_j and with the same relative weights and so replacing individual values with averages in a C_i' would make no difference to the overall average in a C_j. This is why the assumption $I_1 \subset I_2$ is important. Without this assumption the outcomes in a C_i' might be split over different C_j and therefore replacing the individual values with averages in a C_i' would affect the calculation of averages over the C_j.

Figures 1 and 2 present the same reasoning in graphical form. The operation of taking conditional expectations can be thought of as a smoothing or blurring operations that erases the variation in the variable whose expectation is being taken and instead gives us a constant over each chunk corresponding to one realisation of the variables in the information set. The smaller the information set, the bigger the chunks, and hence the more drastic the blurring. Then the law of iterated expectation says that if we want to blur over some big boxes then there is no harm in blurring over some small boxes first provided that each small box lies within a single big box.

1.3. Uses.

1.3.1. *Simplifying calculations.* The first use of the law of iterated expectation is to simplify the calculation of complicated expectations. Examples can be found in [**Ros09**]. Two results which are useful in such calculations are as follows.

PROPOSITION 4.1. *For any two random variables X and Y and information set I, if X belongs to the information set I and $h(X)$ is any function of X then*

$$E[h(X)Y|I] = h(X)E[Y|I].$$

PROPOSITION 4.2. *For any random variable Y and information set I, if Y is independent of the variables in I then*

$$E[Y|I] = E[Y].$$

The two results above correspond to two extremes of information availability. In the first case knowing I gives us full information about X and therefore functions of X behave like constants and come out of expectations. In the second case knowing I gives us no information about Y and therefore the conditional expectation is the same as the unconditional expectations.

1.3.2. *Gradual Revelation of Information.* The past is known whereas the future is unknown. The future, though, depends on the past and therefore knowledge about the past is helpful in predicting the possibilities for the future. As we flow through time, moments pass from
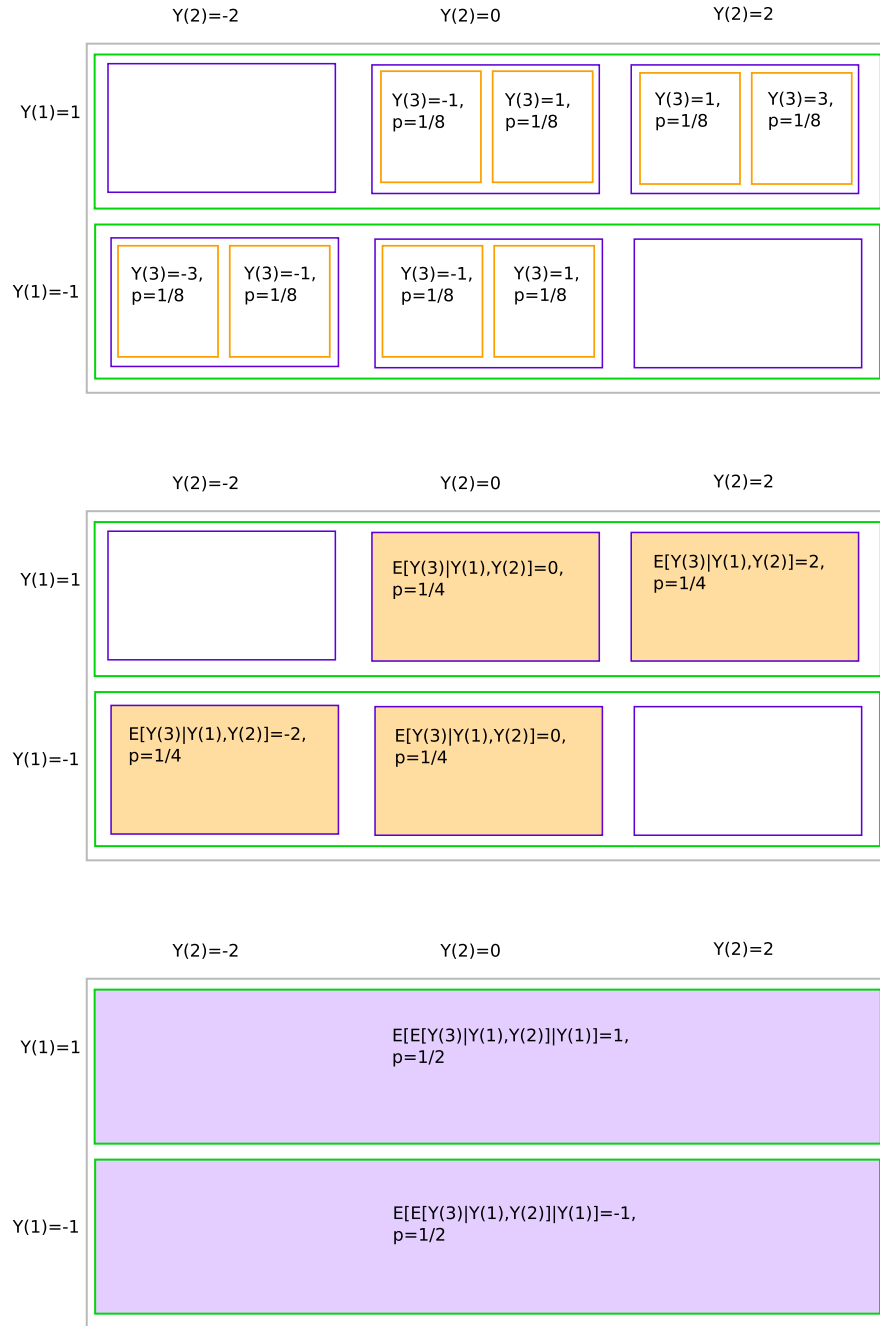
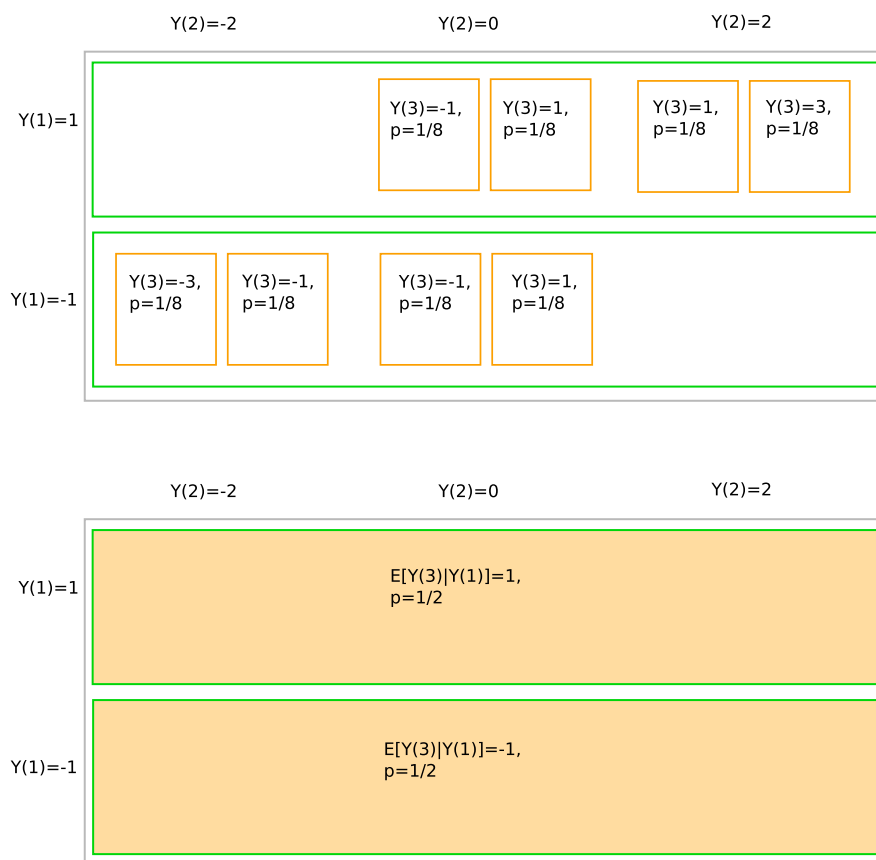FIGURE 1. Law of iterated expectations, left-hand side

FIGURE 2. Law of iterated expectations, right-hand side

the future, through the present, and into the past, resulting in an accumulation of knowledge and a change in our beliefs about the future. A model of macroeconomic dynamics must incorporate these facts about time and knowledge while not giving up on the machinery for talking about uncertainty that the probabilists have developed.

The way we do this is as follows. We set up our probability model in a 'timeless' way. Thus in Example 4.1 even though the three tosses happen over time we think of each point in sample space as specifying the result of all three tosses. For each point of time t an information set I_t is given which specifies what an agent knows at that time. This information set breaks up the sample space into chunks. Thus, for example, if we observe that the first toss has come up heads we know that the outcome must be one of HHH, HHT, HTH or HTT. If the

first toss had come up tails we would have known that the outcome must have been one of THH, THT, TTH or TTT. These are precisely the sets C_1 and C_2 in Table 1. Once the results of the second toss comes in our information set becomes larger and we can narrow down the possible outcomes and know that we are in one of the sets $C'_1, \ldots, C'_2$ in Table 2. Thus the gradual accumulation of knowledge is modelled by adding an increasing sequence of information sets to a timeless probability model.[1]

To be consistent with this formalisation we require that agents choices be constant over each chunk created by their information sets. So if an agent has to make a decision after observing only the first toss then the agent cannot decide differently for the two outcome HHT and HHH since at that time the agent does not have enough information to distinguish between those two outcomes. The agent can however choose differently for the two outcomes HHT and THT since knowing the result of the first toss allows her to discriminate between these two outcomes. This requirement that a random variable like the agent's choice be constant across each chunk created by an information set has a formal name—we say that the random variable is *measurable* with respect to the information set. By construction expectation conditional on an information set are always measurable with respect to that information set.

The notion of measurability allows us to add one more trick useful for working with conditional expectations.

PROPOSITION 4.3. *If I is some information set, Y is any random variable and X is a random variable measurable with respect to I then*

$$E[XY|I] = XE[Y|I]$$

This is a generalisation of Proposition 4.1. If X is measurable with respect to I then it is constant on each chunk created by I and can therefore can be taken out like a constant while calculating expectations conditional on I.

Agents in macroeconomic models must often forecast the value of future random variables based on the limited information that they have at an earlier point of time. The *rational expectations* assumption, which is the most common assumption about expectations currently in macroeconomics, says that this forecast is equal to the conditional

---

[1]Our increasing sequences of information sets is a special case of the more general concept of *filtrations* used by probabilists to model the accumulation of information. See [**Bil95**] for details.

expectation of the future variable based on the agent's current information set. So if $Y_3^{e,1}$ is an agent's forecast after the first toss of final winnings in the game in Example 4.1, then rational expectation says that

$$Y_3^{e,1} = E[Y_3|I_1].$$

The left-hand side of the above equation is an economic agent's subjective forecast while the right-hand side is a mathematical expectation calculated from a probability model. There is no logical reason why they have to be equal, only an economic assumption that they are.

Sometimes we have to go further and make forecasts about forecasts. For example, we may want to forecast after the first toss what we would think of our final winnings after seeing the second toss. Applying the assumption of rational expectations to both these forecasts we get the answer

$$E[E[Y_3|I_2]|I_1].$$

The law of iterated expectations allows us to simplify this to

$$E[Y_3|I_1].$$

Thus the forecast of a forecast is the same as a direct forecast. We shall take advantage of this simplification again and again in our future work in macroeconomics. It is one of the reasons that we give so much importance to the law of iterated expectations.

2. Martingales

Let X_t be a sequence of random variables and I_t be a sequence of information sets such that,

(1) $I_t \subset I_{t+1}$.
(2) X_t is measurable with respect to I_t.
(3) $E[|X_t|] < \infty$.
(4) $E[X_{t+1}|I_t] = X_t$.

Then the process X_t is said to be a *martingale* with respect to the sequence of information sets I_t.

Among the defining characteristics of a martingale given above the crucial one is the last: $E[X_{t+1}|I_t] = X_t$. This says that conditional on the information available in period t the random variable X_t neither grows nor declines in the next period in expected value terms. In this a martingale process is like the wealth of a gambler playing a fair game. Some outcomes of the game increase the gambler's wealth, other outcomes of the game decrease the gambler's wealth, but since the game is fair the increases and decreases cancel out on an average.

We can use Proposition [4.3] and the properties of a martingale to show that

$$E[X_{t+1} - X_t | I_t] = 0$$

and therefore if Y is any variable belonging to I_t then

$$E[(X_{t+1} - X_t)Y] = 0$$

so that $X_{t+1} - X_t$ is uncorrelated with Y. (Give detailed proofs of all the claims made in this paragraph so far.) That is, the change in the value of a martingale process is uncorrelated with all the information available from the past. Speaking loosely, it is a 'surprise'. And so a martingale is a sum of surprises.

A random walk is an martingale (Prove. Be careful to specify the sequence of information sets.) This has led some economic literature to loosely use the term "random walk" when discussing martingales. This usage should be avoided since not all martingales are random walks.

The definition of a martingale tells us about the expectation of a value of the process in a period conditional on the information in the exactly preceding period. The following proposition covers the case where conditioning set and the value is separated by more than one period.

PROPOSITION 4.4. *Let X_t be a martingale with respect to the information set sequence I_t. For any m and any $n > 0$,*

$$E[X_{m+n} | I_m] = X_m.$$

PROOF. The proof is by mathematical induction on n.

For $n = 1$ the proof follows directly from the definition of a martingale.

Suppose the lemma is true for $n = k$. Consider the case $n = k + 1$. From the law of iterated expectations

$$E[X_{m+k+1} | I_m] = E[E[X_{m+k+1} | I_{m+k}] | I_m].$$

The martingale property, applied at time $m + k$ tells us that

$$E[X_{m+k+1} | I_{m+k}] = X_{m+k}.$$

The assumption that the lemma is true for $n = k$ gives us,

$$E[X_{m+k} | I_m] = X_m.$$

Putting everything together, we have

$$E[X_{m+k+1} | I_m] = E[E[X_{m+k+1} | I_{m+k}] | I_m] = E[X_{m+k} | I_m] = X_m$$

thus establishing the result for $n = k + 1$.

Since we have shown that the result is true for $n = 1$ and it is true for $n = k + 1$ whenever it is true for $n = k$, it follows from the principle of mathematical induction that it is true for all $n > 0$. $\square$

CHAPTER 5

Consumption: Uncertainty

1. Euler equation

In the case of uncertainty in labour incomes, but with certain interest rates, the Euler equation becomes

$$v_t'(c_t) = RE_t[v_{t+1}'(c_{t+1})] \tag{32}$$

where E_t denotes the mathematical expectation conditional on information at time t.

2. Quadratic felicity

2.1. Martingale property. Suppose the felicity (i.e. per-period utility) function is

$$v_t(c_t) = \beta^t(ac_t - 0.5c_t^2)$$

where a is some constant.

In this case (32) specialises to

$$a - c_t = R\beta(a - E_tc_{t+1})$$

If we further assume that $R = 1/\beta$ then

$$E_tc_{t+1} = c_t \tag{33}$$

that is, consumption is a martingale process.

Since c_t is part of the information set at time t, $E_tc_t = c_t$ Therefore, another way to write (33) is

$$E_t(c_{t+1} - c_t) = 0$$

which says that the change in consumption between time t and $t + 1$ has no predictable direction based on information at time t.

This result is a consequence of the very special assumptions that we have made. Assuming the same felicity function for each period (apart from the discount factor β) and then assuming that the market rate of discount $(1/R)$ equals this subjective discount factor creates a situation where the consumer has no desire to have a higher consumption in any particular period of her life either to meet greater consumption needs or to take advantage of the difference between market and subjective discount rates. Unconstrained lending and borrowing mean

that the consumer can actually move around her income across periods so as to achieve this perfect symmetry in her consumption in the sense of equating expected marginal utility across periods. But with quadratic felicity expected marginal utility is the same thing as expected consumption and we have our martingale result.

The long list of assumptions leading up to the martingale result means that this precise result is not very robust or realistic. Therefore rather than taking it as a property that is likely to be literally true, we should understand it as a demonstration of the tendency of the lending and borrowing behaviour of consumers to delink current consumption from current income. This tendency will be there as long as consumers have access to asset markets, though in more realistic settings it will be overlaid with factors which impart a systematic pattern to the trajectory of consumption such as a changing pattern of lifetime consumption needs or differences between the subjective and market rate of discount.

2.2. The level of consumption. The martingale property of consumption only tells us how consumption evolves from one point to the next, not the *level* of consumption. The level of consumption would depend on the consumer's resources in terms of her initial wealth and expected labour income. We now show that this is so mathematically by deriving an explicit formula for the level of consumption in the case where consumption is a martingale.

Consider a consumer who stands at period t with wealth w_t and is planning her future consumption for the periods $t, t + 1, \ldots, T$. Since she cannot leave any bequests or outstanding debt in period T, it must be the case that her *realized* stream of consumption (c_t) and labour income (y_t) must satisfy,

$$\sum_{i=t}^{T} \delta^{i-t} c_i = \sum_{i=t}^{T} \delta^{i-t} y_i + w_t \tag{34}$$

Taking expectations as of time t,

$$\sum_{i=t}^{T} \delta^{i-t} E_t c_i = \sum_{i=t}^{T} \delta^{i-t} E_t y_i + w_t \tag{35}$$

From Proposition 4.4 on page 32, $E_t c_i = c_t$ for all $i > t$. And for $i = t$, $E_t c_t = c_t$ because c_t is known at time t. Hence,

$$c_t \sum_{i=t}^{T} \delta^{i-t} = \sum_{i=t}^{T} \delta^{i-t} E_t y_i + w_t \qquad (36)$$

$$c_t = \frac{1}{\sum_{i=t}^{T} \delta^{i-t}} \left[\sum_{i=t}^{T} \delta^{i-t} E_t y_i + w_t \right] \qquad (37)$$

Thus the level of consumption in a given period depends on the expected discounted value of the future stream of labour income over the entire remaining lifetime as well as initial wealth. This once again reiterates the idea of the permanent income hypothesis that the consumption in each period depends not just on income in that period but on the entire expected path of future income.

2.3. Increments in consumption. With an explicit formula for the level of consumption in hand, we can now try to understand the martingale result better by seeing what it is exactly that drives changes in consumption.

Rewriting (36) for period $t + 1$ we have,

$$c_{t+1} \sum_{i=t+1}^{T} \delta^{i-t-1} = \sum_{i=t+1}^{T} \delta^{i-t-1} E_{t+1} y_i + w_{t+1}$$

Substituting $w_{t+1} = (w_t + y_t - c_t)/\delta$,

$$c_{t+1} \sum_{i=t+1}^{T} \delta^{i-t-1} = \sum_{i=t+1}^{T} \delta^{i-t-1} E_{t+1} y_i + (w_t + y_t - c_t)/\delta$$

Multiplying throughout by δ,

$$c_{t+1} \sum_{i=t+1}^{T} \delta^{i-t} = \sum_{i=t+1}^{T} \delta^{i-t} E_{t+1} y_i + w_t + y_t - c_t$$

Subtracting (36) from this equation

$$(c_{t+1} - c_t) \sum_{i=t+1}^{T} \delta^{i-t} - c_t = \sum_{i=t+1}^{T} \delta^{i-t} (E_{t+1} y_i - E_t y_i) + y_t - E_t y_t - c_t$$

But $E_t y_t = y_t$ since income at time t is known at that time,

$$(c_{t+1} - c_t) \sum_{i=t+1}^{T} \delta^{i-t} = \sum_{i=t+1}^{T} \delta^{i-t}(E_{t+1}y_i - E_t y_i)$$

$$c_{t+1} - c_t = \frac{1}{\sum_{i=t+1}^{T} \delta^{i-t}} \left[\sum_{i=t+1}^{T} \delta^{i-t}(E_{t+1}y_i - E_t y_i) \right]$$

We can divide both the numerator and denominator by δ to have the factor multiplying the first term in the sums equal to one.[1] This gives us our final formula,

$$c_{t+1} - c_t = \frac{1}{\sum_{i=t+1}^{T} \delta^{i-t-1}} \left[\sum_{i=t+1}^{T} \delta^{i-t-1}(E_{t+1}y_i - E_t y_i) \right] \qquad (38)$$

What the formula above says is that changes in consumption are a result of *revisions* in expectations of future income based on the difference in information at time t and $t + 1$. Therefore changes in income that were predictable at time t do not contribute to the change in consumption between time t and $t + 1$. This is consistent with our assumption that consumption is a martingale but goes further by predicting the actual size of the consumption change rather than just asserting that the expected value of this change would be zero.

2.4. Specific income processes. In general the difference in expectations which occur on the right of (38) depends on all the information which becomes available to the consumer between time t and $t+1$. One special case which we now consider is when the only source of new information is the realisation of the labour income y_{t+1}. What revision this new information causes in the expectation of future labour income depends on how labour income in different periods are related.

As an example consider the labour income process given by the following stochastic difference equation,

$$y_{t+1} - \mu = \rho(y_t - \mu) + \epsilon_{t+1} \qquad (39)$$

where ϵ_t is a white-noise process, μ and ρ are constants and $0 < \rho < 1$. This is a special case of what is known as a first-order autoregressive process (sometimes denoted as a AR(1) process). The coefficient ρ measures how persistent the deviations in y from its long-run average μ are.

---

[1]This is a purely aesthetic change and does not change any results.

Writing (39) for the period $t + 2$ we have,

$$y_{t+2} - \mu = \rho(y_{t+1} - \mu) + \epsilon_{t+2}$$

Substituting (39),

$$y_{t+2} - \mu = \rho^2(y_t - \mu) + \rho\epsilon_{t+1} + \epsilon_{t+2}$$

Carrying out successive substitutions like this, we find for any $i > t$

$$y_i - \mu = \rho^{i-t}(y_t - \mu) + \sum_{j=t+1}^{i} \rho^{i-j}\epsilon_j$$

Taking expectations conditional on the information at time t,

$$E_t(y_i - \mu) = \rho^{i-t}(y_t - \mu) + \sum_{j=t+1}^{i} \rho^{i-j}E_t\epsilon_j$$

Here we have used the fact that y_t is known at time t. We further note that since ϵ_t is IID, ϵ_j is independent of all information at time t when $j > t$ and we can replace $E_t\epsilon_j$ by $E\epsilon_j$ which is 0 by the definition of white noise. Hence we conclude,

$$E_t(y_i - \mu) = \rho^{i-t}(y_t - \mu) \qquad \text{for } i \geq t \tag{40}$$

(We have established this above for $i > t$ and it is trivially true for $i = t$.)

Using $t + 1$ in the place of t,

$$E_{t+1}(y_i - \mu) = \rho^{i-t-1}(y_{t+1} - \mu) \qquad \text{for } i \geq t + 1 \tag{41}$$

For $i \geq t + 1$ both (40) and (41) hold. Subtracting the former from the latter we have,

$$E_{t+1}y_i - E_t y_i = \rho^{i-t-1}(y_{t+1} - \mu) - \rho^{i-t}(y_t - \mu)$$

Using (39),

$$= \rho^{i-t-1}[\rho(y_t - \mu) + \epsilon_{t+1}] - \rho^{i-t}(y_t - \mu)$$
$$= \rho^{i-t-1}\epsilon_{t+1}$$

Substituting this in (38) we get,

$$c_{t+1} - c_t = \left[\frac{\sum_{i=t+1}^{T}(\delta\rho)^{i-t-1}}{\sum_{i=t+1}^{T} \delta^{i-t-1}} \right] \epsilon_{t+1}$$

So we see that for a given innovation in consumption, ϵ_{t+1}, the increment in consumption is higher the higher is the degree of persistence ρ in the income process.

In empirical applications of the model we can estimate ρ (or its analogues for more complex income processes) from data on consumers' incomes and then check if changes in consumption satisfy the forumla above. This yields a sharper test of our theory compared to just checking if consumption is a martingale.

Bibliography

[Bil95] Patrick Billingsley. *Probability and Measure*. Wiley-Interscience, 3rd edition, 1995.

[CA03] K. L. Chung and Farid AitSahlia. *Elementary Probability Theory: With Stochastic Processes and an Introduction to Mathematical Finance.* Springer, 4th edition, 2003.

[JP04] Jean Jacod and Philip Protter. *Probability Essentials*. Springer, 2nd edition, 2004.

[Ros09] Sheldon M. Ross. *Introduction to Probability Models*. Academic Press, 10th edition, 2009.